

Data exploration with Pandas

CS424: Visualization & Visual Analytics

Fabio Miranda

<https://fmiranda.me>

Pandas



- Powerful Python package for manipulating tables.
- Built on top of numpy.
- Save time by abstracting lower-level code for manipulating, extracting, and deriving data tables.
- Easy & quick visualization with matplotlib.
- Main data structures: **Series** and **DataFrame**

Simple series

```
1 data = pd.Series([0.1, 0.2, 0.3, 0.4, 0.5])
```

```
1 data
```

0	0.1
1	0.2
2	0.3
3	0.4
4	0.5

dtype: float64

Index
Data

Explicit index

```
1 data = pd.Series([0.1, 0.2, 0.3, 0.4, 0.5], index = ['a', 'b', 'c', 'd', 'e'])
```

```
1 data
```

a	0.1
b	0.2
c	0.3
d	0.4
e	0.5

dtype: float64

Indexing & accessing data

- `.loc` label based:
 - A single label
 - A list of array of labels
 - A slice object with labels
 - A boolean array
 - A callable function with one argument that returns valid output for indexing (one of the above).
- `.iloc` integer position based:
 - An integer
 - A list of array of integers
 - A slice object with integers
 - A boolean array
 - A callable function with one argument that returns valid output for indexing (one of the above)

Indexing & accessing data

Selection using label

```
1 data.loc['a'] # single label
```

```
0.1
```

```
1 data.loc[['a','b']] # list of labels
```

```
a    0.1  
b    0.2  
dtype: float64
```

```
1 data.loc['a':'c'] # slice object with labels
```

```
a    0.1  
b    0.2  
c    0.3  
dtype: float64
```

```
1 data.loc[[False,False,True,False,False]] # boolean mask
```

```
c    0.3  
dtype: float64
```

```
1 data.loc[lambda x: x.index == 'b'] # callable function
```

```
b    0.2  
dtype: float64
```

Selection using integer position

```
1 data.iloc[0] # scalar integer
```

```
0.1
```

```
1 data.iloc[[0,1]] # list of integers
```

```
a    0.1  
b    0.2  
dtype: float64
```

```
1 data.iloc[0:2] # slice object
```

```
a    0.1  
b    0.2  
dtype: float64
```

```
1 data.iloc[[False,False,True,False,False]] # boolean mask
```

```
c    0.3  
dtype: float64
```

```
1 data.iloc[lambda x: x.index == 'b'] # callable function
```

```
b    0.2  
dtype: float64
```

Dictionary as a series

```
1 population_dict = {'California': 38332521,  
2                    'Texas': 26448193,  
3                    'New York': 19651127,  
4                    'Florida': 19552860,  
5                    'Illinois': 12882135}
```

```
1 population = pd.Series(population_dict)  
2 population
```

```
California    38332521  
Texas         26448193  
New York      19651127  
Florida       19552860  
Illinois      12882135  
dtype: int64
```

```
1 population.loc['California']
```

```
38332521
```

```
1 population.loc[population > 20000000]
```

```
California    38332521  
Texas         26448193  
dtype: int64
```

Accessing with boolean array

DataFrame object

- DataFrame is a 2-dimensional labeled data structure with columns of (potentially) different types.
 - Just like a spreadsheet or SQL table, or dict of Series objects.
- DataFrame can be created with:
 - Dict of 1D arrays, lists, dicts, or Series
 - 2D numpy array
 - Series
 - Another DataFrame

Constructing a DataFrame

- From a dictionary or list of dictionaries:

```
1 d = {"one": [1.0, 2.0, 3.0, 4.0]}
2 pd.DataFrame(d)
```

	one
0	1.0
1	2.0
2	3.0
3	4.0

```
1 d = {"one": [1.0, 2.0, 3.0, 4.0], "two": [4.0, 3.0, 2.0, 1.0]}
2 pd.DataFrame(d)
```

	one	two
0	1.0	4.0
1	2.0	3.0
2	3.0	2.0
3	4.0	1.0

```
1 d = {"one": [1.0, 2.0, 3.0, 4.0], "two": [4.0, 3.0, 2.0, 1.0]}
2 pd.DataFrame(d, index=["a", "b", "c", "d"])
```

	one	two
a	1.0	4.0
b	2.0	3.0
c	3.0	2.0
d	4.0	1.0

- From numpy ndarray:

```
1 pd.DataFrame(np.random.randint(low=0, high=10, size=(5,5)), columns=['a', 'b', 'c', 'd', 'e'])
```

	a	b	c	d	e
0	8	4	6	1	1
1	1	8	3	8	8
2	2	7	9	2	1
3	5	8	4	9	3
4	0	0	6	9	8

Constructing a DataFrame

- From dictionaries or Series

```
1 population_dict = {'California': 38332521,  
2                   'Texas': 26448193,  
3                   'New York': 19651127,  
4                   'Florida': 19552860,  
5                   'Illinois': 12882135}
```

```
1 area_dict = {'California': 423967,  
2             'Texas': 695662,  
3             'New York': 141297,  
4             'Florida': 170312,  
5             'Illinois': 149995}
```

```
1 states = pd.DataFrame({'population': population_dict, 'area': area_dict})  
2 states
```

	population	area
California	38332521	423967
Texas	26448193	695662
New York	19651127	141297
Florida	19552860	170312
Illinois	12882135	149995

Viewing data & statistics

```
1 states.head(2)
```

	population	area
California	38332521	423967
Texas	26448193	695662

```
1 states.tail(2)
```

	population	area
Florida	19552860	170312
Illinois	12882135	149995

```
1 states.describe()
```

	population	area
count	5.000000e+00	5.000000
mean	2.337337e+07	316246.600000
std	9.640386e+06	242437.411951
min	1.288214e+07	141297.000000
25%	1.955286e+07	149995.000000
50%	1.965113e+07	170312.000000
75%	2.644819e+07	423967.000000
max	3.833252e+07	695662.000000

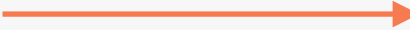
Computing descriptive stats
for each column

Viewing sorted DataFrame

```
1 states
```

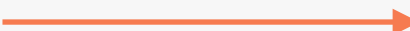
	population	area
California	38332521	423967
Texas	26448193	695662
New York	19651127	141297
Florida	19552860	170312
Illinois	12882135	149995

```
1 states.sort_index()
```

 Viewing sorted by index

	population	area
California	38332521	423967
Florida	19552860	170312
Illinois	12882135	149995
New York	19651127	141297
Texas	26448193	695662

```
1 states.sort_values(by='area')
```

 Viewing sorted by column

	population	area
New York	19651127	141297
Illinois	12882135	149995
Florida	19552860	170312
California	38332521	423967
Texas	26448193	695662

Selecting & filtering data

- Selection using integer position
- Multi-axis selection by label

```
1 states.iloc[0]
```

```
population    38332521  
area          423967  
Name: California, dtype: int64
```

```
1 states.loc[:, ['population']]
```

population	
California	38332521
Texas	26448193
New York	19651127
Florida	19552860
Illinois	12882135

```
1 states.loc[['New York', 'Illinois'], ['population']]
```

population	
New York	19651127
Illinois	12882135

Selecting & filtering data

- Boolean indexing

```
1 states[states['population'] > 20000000]
```

	population	area
California	38332521	423967
Texas	26448193	695662

```
1 states[states.index.isin(['New York'])]
```

	population	area
New York	19651127	141297

Operations

```
1 d = pd.DataFrame(np.random.randint(low=0, high=10, size=(5,5)), columns=['a', 'b', 'c', 'd', 'e'])
```

```
1 d
```

	a	b	c	d	e
0	2	9	7	7	7
1	5	8	3	7	3
2	8	3	0	0	1
3	1	9	8	0	0
4	4	0	3	9	2

```
1 d.mean()
```

→ Across axis 0 (rows), i.e., column mean

```
a    4.0
b    5.8
c    4.2
d    4.6
e    2.6
dtype: float64
```

```
1 d.mean(axis=1)
```

→ Across axis 1 (columns), i.e., row mean

```
0    6.4
1    5.2
2    2.4
3    3.6
4    3.6
dtype: float64
```

Operations

```
1 d.apply(np.cumsum) → NumPy's cumulative sum
```

	a	b	c	d	e
0	2	9	7	7	7
1	7	17	10	14	10
2	15	20	10	14	11
3	16	29	18	14	11
4	20	29	21	23	13

```
1 states.apply(lambda x: x['population'] / x['area'], axis=1)
```

```
California    90.413926
Texas         38.018740
New York      139.076746
Florida       114.806121
Illinois      85.883763
dtype: float64
```

Population density of each **row**

Merging tables

```
1 left = pd.DataFrame({'key': ['foo', 'bar'], 'lval': [1,2]})  
2 right = pd.DataFrame({'key': ['foo', 'bar'], 'lval': [4,5]})
```

```
1 left
```

	key	lval
0	foo	1
1	bar	2

```
1 right
```

	key	lval
0	foo	4
1	bar	5

```
1 pd.merge(left, right, on='key') → Column or index names to join on
```

	key	lval_x	lval_y
0	foo	1	4
1	bar	2	5

Grouping

```
1 df = pd.DataFrame({'Animal': ['Falcon', 'Falcon',  
2                           'Parrot', 'Parrot'],  
3                     'Max Speed': [380., 370., 24., 26.]})
```

```
1 df
```

	Animal	Max Speed
0	Falcon	380.0
1	Falcon	370.0
2	Parrot	24.0
3	Parrot	26.0

```
1 df.groupby(['Animal']).mean()
```

	Max Speed
Animal	
Falcon	375.0
Parrot	25.0

Grouping

```
1 arrays = [['Falcon', 'Falcon', 'Parrot', 'Parrot'],
2           ['Captive', 'Wild', 'Captive', 'Wild']]
3 index = pd.MultiIndex.from_arrays(arrays, names=('Animal', 'Type'))
4 df = pd.DataFrame({'Max Speed': [390., 350., 30., 20.]}, index=index)
```

```
1 df
```

Max Speed		
Animal	Type	
Falcon	Captive	390.0
	Wild	350.0
Parrot	Captive	30.0
	Wild	20.0

```
1 df.index
```

```
MultiIndex([('Falcon', 'Captive'),
            ('Falcon', 'Wild'),
            ('Parrot', 'Captive'),
            ('Parrot', 'Wild')],
            names=['Animal', 'Type'])
```

Grouping by index:

```
1 df.groupby(level=0).mean()
```

Max Speed	
Animal	
Falcon	370.0
Parrot	25.0

```
1 df.groupby(level="Type").mean()
```

Max Speed	
Type	
Captive	210.0
Wild	185.0

Importing & exporting data

- Reading and writing a CSV file:

```
1 pd.read_csv('data.csv')
```

```
1 df.to_csv('data.csv')
```

- DataFrame to binary Feather format:

```
1 df.to_feather('data.feather')
```

Basic plotting with matplotlib

```
1 ts = pd.Series(np.random.randn(1000), index=pd.date_range("1/1/2000", periods=1000))
2 df = pd.DataFrame(np.random.randn(1000, 4), index=ts.index, columns=list("ABCD"))
3 df = df.cumsum()
4 df
```

	A	B	C	D
2000-01-01	0.099142	-0.679263	-0.669535	0.971732
2000-01-02	-0.713262	-1.037180	-1.869124	0.314566
2000-01-03	-2.176599	-2.202236	-0.843755	-0.426149
2000-01-04	-1.254498	-2.075695	-2.420534	0.228423
2000-01-05	-0.251042	0.105400	-2.590070	0.277761
...
2002-09-22	11.209192	24.387028	27.601228	-87.805667
2002-09-23	12.023897	23.530602	26.630084	-88.124066
2002-09-24	10.766121	23.579338	26.731239	-87.990660
2002-09-25	11.518224	23.913193	27.140907	-86.354709
2002-09-26	12.567776	24.353585	27.994359	-86.652313


1000 rows × 4 columns

Basic plotting with matplotlib

```
1 plt.figure()  
2 df.plot()
```

<AxesSubplot:>


<Figure size 432x288 with 0 Axes>



Basic plotting with matplotlib


```
1 df.loc['1/10/2000'].plot()
```

<AxesSubplot:>



```
1 df.iloc[9].plot()
```

<AxesSubplot:>




Basic plotting with matplotlib

- *bar* for bar plots
- *hist* for histogram
- *box* for boxplot
- *kde* for density plots
- *area* for area plots
- *scatter* for scatter plots
- ...

```
1 df.iloc[9].plot(kind='bar')
```


<AxesSubplot:>



Basic plotting with matplotlib


```
1 df = pd.DataFrame(np.random.rand(10, 4), columns=["a", "b", "c", "d"])
2 df.plot(kind='bar')
```

<AxesSubplot:>




```
1 df.plot(kind='bar', stacked=True)
```

<AxesSubplot:>



```
1 df.plot(kind='barh', stacked=True)
```

<AxesSubplot:>



```
1 df = pd.DataFrame(np.random.rand(50, 4), columns=["a", "b", "c", "d"])
2 df.plot.scatter(x="a", y="b")
```

