

# Spatial Data Structures

CS425: Computer Graphics I

Fabio Miranda

<https://fmiranda.me>


# Overview

---

- Spatial data structures
  - Uniform grid
  - Nested grids
  - Quadtree / octree
  - K-d tree
  - Bounding Volume Hierarchy
  - Efficient sparse voxel octrees
  - High resolution sparse voxel DAGs

# How to efficiently organize objects?

- 3D data contains spatial information.
- How to perform queries when there are thousands / millions of objects (points, polygons)?
  - Ray-scene intersection.
  - Proximity queries
  - Point in polygon.
  - Range query.




# Ray-scene intersection

- Given a scene with  $n$  primitives and a ray  $r$ , find the closest point of intersection of  $r$  with the scene.

```
function intersectObjects(ray, scene) {  
  for(var i=0; i < scene.objects.length; i++) {  
    var object = scene.objects[i];  
    var dist = intersection(ray, object);  
    // ...  
  }  
}
```

- Complexity:  $O(n)$
- How to do better?

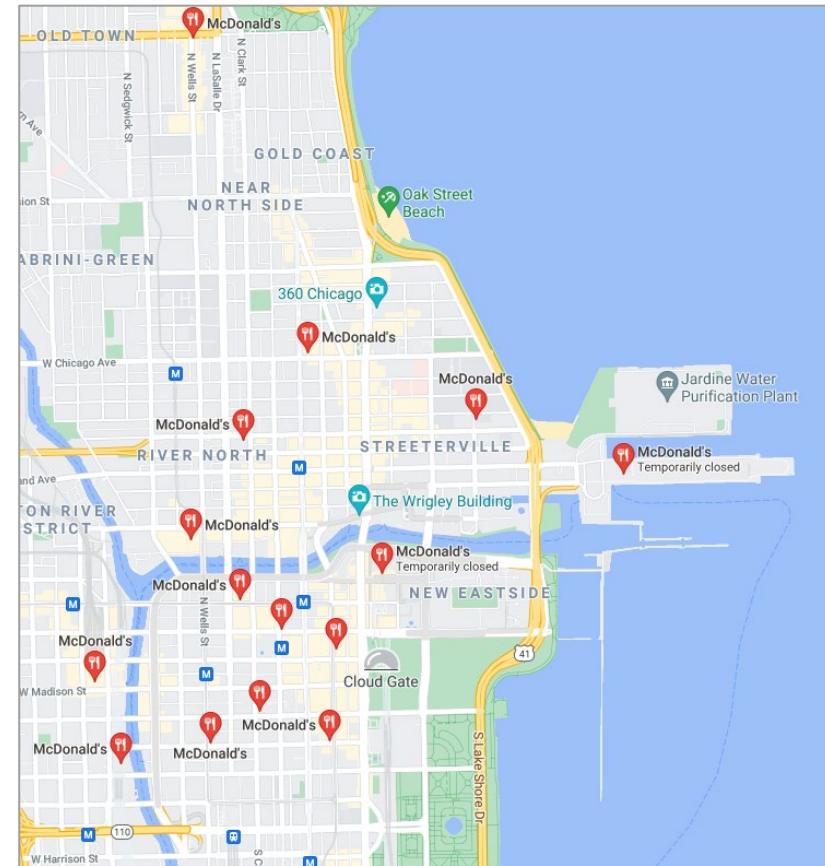


# Proximity query

- Query based on proximity.
- “What is the closest McDonald’s?”

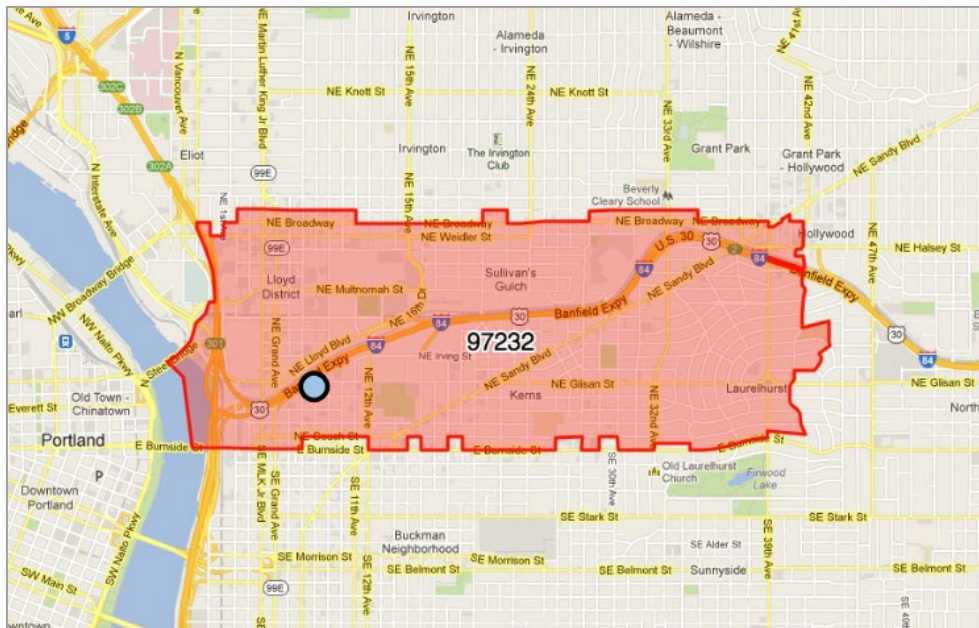
```
function findPlaces(query, scene) {  
    for(var i=0; i < scene.places.length; i++) {  
        var place = scene.places[i];  
        var dist = satisfyQuery(query, place);  
        // ...  
    }  
}
```

- Complexity:  $O(n)$
- How to do better?

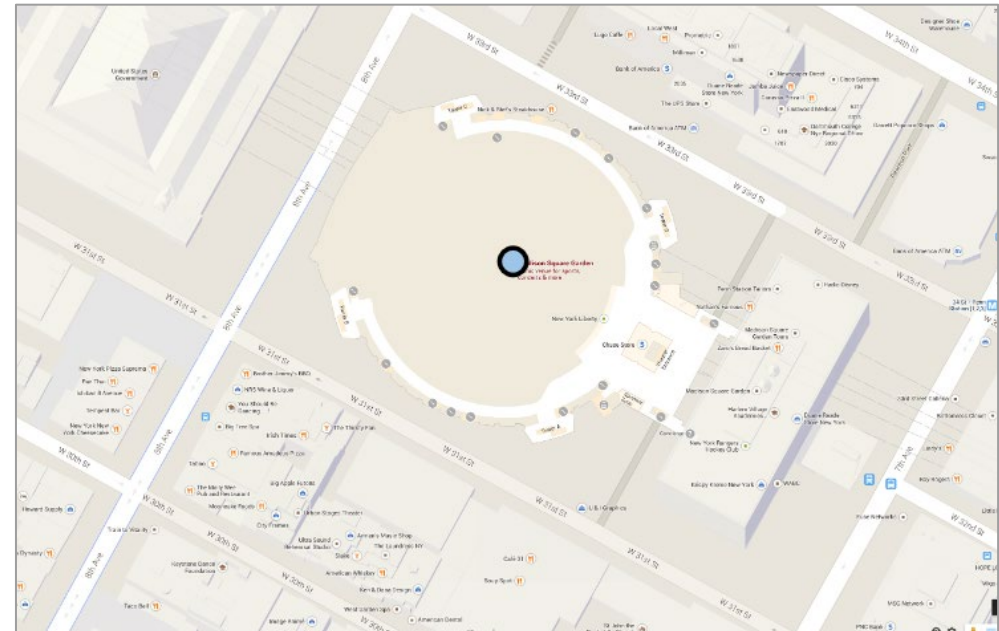




# Point in polygon



What is the zip code for this complaint?



Am I inside a specific building?

# Time complexity

- Ray-scene intersection:  $O(n)$
- Proximity query:  $O(n)$
- Point in polygon:  $O(n)$

**How to reduce the time complexity?**

# Motivation

---

- Expensive operations (ray tracing, query).
  - Complex scenes (millions of objects).
  - Large number of operations (hundreds of millions per second).
- Reduce complexity through pre-processing data.
  - Spatial data structures: structures of objects in space.
  - Eliminate candidates as early as possible.
  - Reduce complexity to  $O(\log n)$  on average.
  - Worst case complexity still  $O(n)$ .
    - Can you come up with a worst case example?



# Spatial data structures


- Data structures to accelerate queries of the kind:  
**“I’m here, which object is around me?”**
- Partition space or set of objects.
- Tasks:
  1. Construction / update:
    - Pre-processing for static parts of the scene.
    - Update for moving parts of the scene.
  2. Access:
    - Optimize so it is done as fast as possible.

# Spatial data structures



- Uniform grid: 2D/3D data, uniform distribution.
- Quadtree: 2D data, non-uniform distribution.
- Octree: 3D data, non-uniform distribution.
- KD-tree: 2D/3D data, avoid empty cells.

# Uniform grid

- Partition space into equal-sized volumes (i.e., voxels).
- Each cell will contain objects that overlap the voxel.
- Good for uniform data (points are evenly distributed in space).
- Fast construction and queries.



# Uniform grid



# Uniform grid: construction and query


- Array of 3D voxels
  - Each voxel: list of pointers to colliding objects.
- Indexing function:
  - 3D point  $\rightarrow$  cell index (constant time!)
- Construction:
  - Initialize cells for grid with size  $w * h$
  - For each object  $p(x, y)$ :
    - Compute grid cell using  $(x, y)$ .
    - Store  $p$  in cell.
- Query:
  - For query rectangle  $(x_1, y_1) \times (x_2, y_2)$ :
    - Compute subgrid for  $(x_1, y_1)$  and  $(x_2, y_2)$ .
    - For all cells inside subgrid, report all objects.
    - For all cells on the border of the subgrid, test objects against rectangle.

# Uniform grid: complexity

- Build time:  $O(n)$
- Space:  $O(w * h) + O(n)$
- Query:  $O(k)$

# Uniform grid: complexity

- When uniform grids work well? Uniform distribution of objects.



Mitsuba renderer






peterguthrie.net



# Uniform grid: drawbacks

- When uniform grids do not perform well? Non-uniform distribution of objects.
- “Teapot in a stadium” problem: uniform grids cannot adapt to local density of objects.





# Uniform grid: drawbacks

- Assumes objects uniformly distributed in space.
- What happens when assumption does not hold?
  - Many empty cells.
  - Few cells with too many points.
- Change cell size?
  - Too small: memory occupancy too large.
  - Too big: too many objects in one cell.

# Nested grids


- Possible solution to “teapot in a stadium” problem.
- Hierarchy of uniform grids: each cell is itself a grid.
- Fast building & traversal.



Philipp Slusallek



# Incremental traversal

- How to traverse the voxels?
- Incremental traversal similar to line rasterization:
  - From current intersection  $P$  on cell  $(i, j)$ , perform ray-plane intersection tests with the next 2 grid planes along the ray direction.
  - Next intersection point is the nearest one among the 2 candidates.
  - Update cell index according to new intersection point.
  - Repeat process until the ray intersects with a primitive or reaches the boundary of the grid.





# Quadtree

- Hierarchical structure that stores regular grids at each level.
- Adaptive subdivision: adjust depth to local scene complexity.




# Quadtree

- Rooted tree in which every internal node has four children.
- Every node corresponds to a square.
- Tree: branching factor 4 or 8.
- Each node: splits into all dimensions at once (in the middle).
- Construction: continue splitting until end nodes have few objects (or limit level reached).



# Quadtree: construction


- Split the top level.






# Quadtree: construction

- Split the top level.
- Can we stop?




# Quadtree: construction

- Split the top level.
- Can we stop? No, split the next level.




# Quadtree: construction

- Split the top level.
- Can we stop? No, split the next level.
- Split top-left.




# Quadtree: construction

- Split the top level.
- Can we stop? No, split the next level.
- Split top-left.
- Can we stop top-left?




# Quadtree: construction

- Split the top level.
- Can we stop? No, split the next level.
- Split top-left.
- Can we stop top-left? Yes.




# Quadtree: construction

- Split the top level.
- Can we stop? No, split the next level.
- Split top-left.
- Can we stop top-left? Yes.
- Split top-right.



# Quadtree: construction


- Split the top level.
- Can we stop? No, split the next level.
- Split top-left.
- Can we stop top-left? Yes.
- Split top-right.
- Can we stop top-right?






# Quadtree: construction

- Split the top level.
- Can we stop? No, split the next level.
- Split top-left.
- Can we stop top-left? Yes.
- Split top-right.
- Can we stop top-right? No.




# Quadtree: construction

- Split the top level.
- Can we stop? No, split the next level.
- Split top-left.
- Can we stop top-left? Yes.
- Split top-right.
- Can we stop top-right? No.
- Split top-right.




# Quadtree: construction

- Split the top level.
- Can we stop? No, split the next level.
- Split top-left.
- Can we stop top-left? Yes.
- Split top-right.
- Can we stop top-right? No.
- Split top-right.
- Can we stop top-right?




# Quadtree: construction

- Split the top level.
- Can we stop? No, split the next level.
- Split top-left.
- Can we stop top-left? Yes.
- Split top-right.
- Can we stop top-right? No.
- Split top-right.
- Can we stop top-right? Yes.




# Quadtree: construction

- Split the top level.
- Can we stop? No, split the next level.
- Split top-left.
- Can we stop top-left? Yes.
- Split top-right.
- Can we stop top-right? No.
- Split top-right.
- Can we stop top-right? Yes.



# Quadtree: construction



# Quadtree: construction

- Construction:
  - Input: set of objects  $P$  inside a square  $S (x_1, y_1) \times (x_2, y_2)$ , tree node  $v$
  - If  $|P| \leq 1$ :
    - Quadtree consists of a single leaf with  $P$ .
  - Else:
    - $P_{00}$ : set of points that fall in the bottom-left corner of  $S$ .
    - $P_{01}$ : set of points that fall in the bottom-right corner of  $S$ .
    - ...
    - $v_{00}$ : node with points of  $P_{00}$ .
    - $v_{01}$ : node with points of  $P_{01}$ .
    - ...
    - Append  $v_{00}, v_{01}, v_{10}, v_{11}$  to  $v$ .



# Quadtree: query


- Query:
  - Input: range query  $r$   $(x_1, y_1) \times (x_2, y_2)$ , tree node  $v$ .
  - If  $v$  is a leaf:
    - Search points of  $v$  inside range  $r$ .
  - If  $v_{00}$  inside range  $r$  :
    - Query( $v_{00}$ ,  $r$ )
  - If  $v_{01}$  inside range  $r$  :
    - Query( $v_{01}$ ,  $r$ )
  - If  $v_{10}$  inside range  $r$  :
    - Query( $v_{10}$ ,  $r$ )
  - If  $v_{11}$  inside range  $r$  :
    - Query( $v_{11}$ ,  $r$ )

# Quadtree: complexity

- Build time:  $O(n)$
- Space:  $O(n)$
- Range query:  $O(\sqrt{n} + k)$
- Leaf traversal:  $O(\log n)$


# Octree

- Each inner node contains 8 equally sized voxels.
- A 3D quadtree.



# Quadtree and octree: drawbacks

- Greater ability to adapt to location of scene geometry than uniform grid.
- But very long tree to store points that are concentrated in a small region.
- Many nodes will contain zero objects.




# K-d tree

---

- Differently from quadtrees and octrees, k-d trees only split one dimension at each level.
- Where to split? Middle? Median? Proportional to surface area?
- At each level:
  - Quadtree creates 4 equal sized cells.
  - Octree creates 8 equal sized cells.
  - K-d tree creates 2 non-equal sized cells (2D case).


# K-d tree: construction

- First split: x dimension (median point).




# K-d tree: construction

- First split: x dimension (median point).
- Second split: y dimension.



# K-d tree: construction

- First split: x dimension (median point).
- Second split: y dimension.
- Repeat, alternating split dimensions





# K-d tree: construction

- Construction:
  - Input: set of objects  $P$  inside a square  $S (x_1, y_1) \times (x_2, y_2)$ , tree node  $v$
  - If  $|P| \leq 1$ :
    - K-d tree consists of a single leaf with  $P$ .
  - Else:
    - If depth is even:
      - Split  $P$  into  $P_0$  and  $P_1$ , along a vertical line through the  $y$  axis.
    - Else:
      - Split  $P$  into  $P_0$  and  $P_1$ , along a vertical line through the  $x$  axis.
    - $v_0$ :  $build(v, P_0, depth + 1)$ .
    - $v_1$ :  $build(v, P_1, depth + 1)$ .
    - ...
    - Append  $v_0, v_1$  to  $v$ .

# K-d tree: query

- Query:
  - Input: range query  $r$   $(x_1, y_1) \times (x_2, y_2)$ , tree node  $v$ .
  - If  $v$  is a leaf:
    - Search points of  $v$  inside range  $r$ .
  - If  $v_0$  inside range  $r$  :
    - Query( $v_0$ ,  $r$ )
  - If  $v_1$  inside range  $r$  :
    - Query( $v_1$ ,  $r$ )


# K-d tree: complexity

- Build time:  $O(n \log n)$
- Space:  $O(n)$
- Range query:  $O(\sqrt{n} + k)$
- Leaf traversal:  $O(\log n)$



# Bounding volume hierarchies (BVH)

- Goal: use the scene hierarchy given by the scene graph (instead of a spatial derived one).
- Associate a bounding volume to each node.
  - A bounding volume of a node bounds **all** objects in the subtree.
  - Nodes are aggregated objects.
- Construction and update is fast.
  - Bottom-up: recursive.
- Querying it:
  - Top-down: visit.



# Bounding volume hierarchies (BVH)



# Bounding volume hierarchies (BVH)



# Bounding volume hierarchies (BVH)



- Leaf nodes: contain *small* list of primitives.
- Interior nodes:
  - Proxy for a large subset of primitives.
  - Stores bounding box for all primitives in subtree.


# Bounding volume hierarchies (BVH)

- How to split the scene?
- Primitive list of current node  $P$  is sorted based on the centroids of primitive bounding boxes. Ordered list is split into  $P_0$  and  $P_1$ .
- Different heuristics:
  - Edge volume heuristic (EVH).
  - Split bounding volume hierarchy (SBVH).
  - Surface area heuristic (SAH):

$$C = C_t + \frac{SA(B_1)}{SA(B)} |P_1| C_i + \frac{SA(B_2)}{SA(B)} |P_2| C_i$$



# Bounding volume hierarchies (BVH)



**Figure 1:** Sample scene consisting of roughly 1.9 million triangles (left). Our method (middle) results in a significant reduction of ray shooting costs compared to a regular bounding volume hierarchy (right). The heat views visualize the summed number of traversal steps and primitive intersections for primary rays.

[Stich et al., 2009]

# Object vs. space partitioning

- Object partitioning:
  - Hierarchically partition objects into groups.
  - Spatial index is created by spatially bounding each subgroup.
- Space partitioning:
  - Hierarchically partition space into subspaces.
  - Subspaces are non-overlapping and completely fill parent space.
  - Tree or table structure.

# Summary

---

- Choose the right structure considering the operations and data.
- Uniform grid:
  - The most parallelizable (to update, construct, use).
  - Constant time access (best!).
  - Quadratic / cubic space (2D, 3D).
  - Good performance under uniform distribution of objects.
- Quadtree, octree, k-d tree:
  - Compact.
  - Simple.
  - Non-constant accessing time.
  - Good performance under non-uniform distribution of objects.
- BVH:
  - Simple construction.
  - Ideal for dynamic parts of the scene.
  - Not necessarily very efficient to access.
    - May need to traverse multiple children.
    - If you do not have a scene-graph you need to create one.

# Voxel-based rendering

- Outcast, 1999.
- *“At this point Yves started wondering if dedicated hardware would ever catch-up the exponential acceleration of CPU power and flexibility”.*
- Flexibility to program effects not available on GPUs at the time:
  - Ripple effects
  - Real-time shadows
  - Curved surfaces




<http://francksauer.com/index.php/games/test/15-games/published-games/47-outcast-pc>

# Sparse voxel octrees

- Using voxel representations as a generic way for expressing complex and feature-rich geometry.
- Voxels as a representation of opaque surfaces.
  - Fine-grained resolution control
  - Compact representation
  - Easy to downsample
- Efficient sparse voxel octrees avoid building the octree with maximum depth by providing each voxel with a contour.
  - If contours are good approximation to the original geometry, subdivision is stopped.

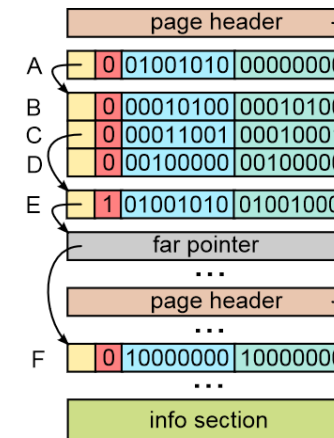


# Sparse voxel octrees: contours

- Contours are slabs that modify voxel's shape.
- Improve accuracy for representing surfaces.
- Simple for rendering.



# Efficient sparse voxel octrees


- How to efficiently store octrees?
  - One child pointer per voxel.
  - Efficient encoding of empty regions.
  - Culling away nodes that correspond to empty space.



Laine and Karras, 2010



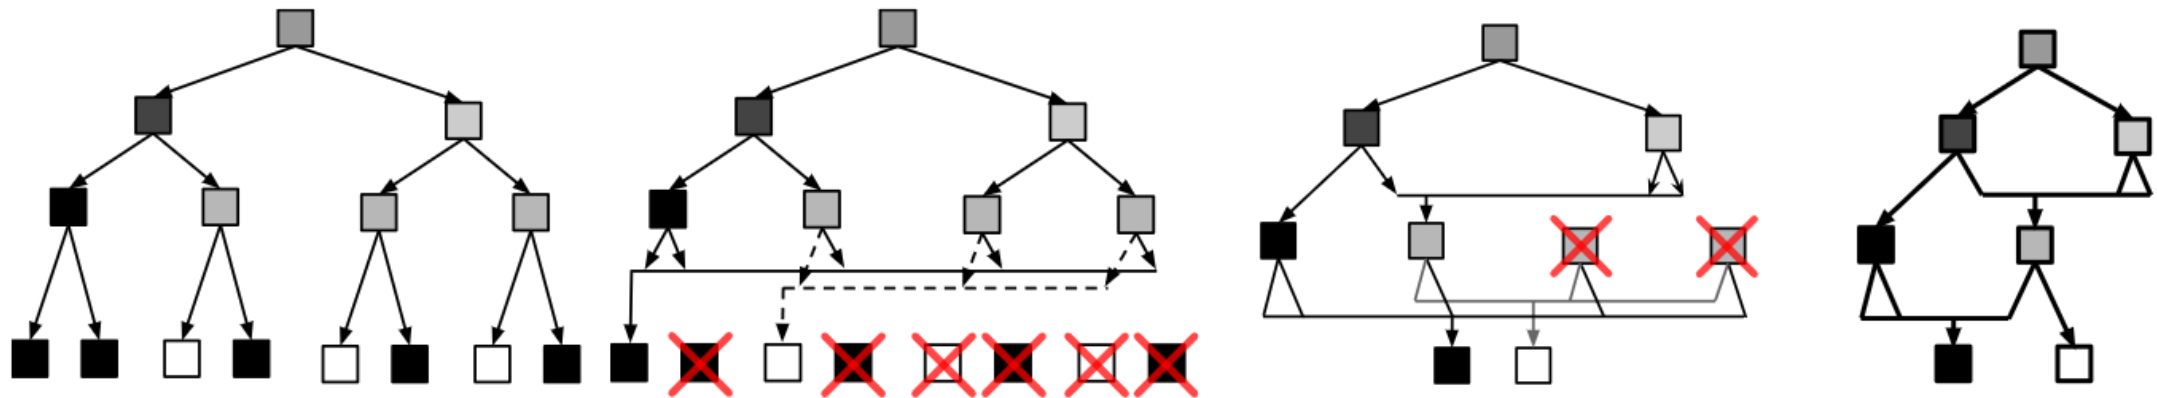
# Efficient sparse voxel octrees





# High resolution sparse voxel DAGs


- Binary voxel grid – generalization of octree to a directed acyclic graph.
- Assumption that SVO contains redundant subtrees.
  - Searches the tree for common subtrees and only reference unique instances.
- DAG representation decreases memory consumption by up to 38x.



# High resolution sparse voxel DAGs

- SVO to DAG:
  - Top down approach:
    - Start at the root node and test whether its children correspond to identical subtrees, and proceed down recursively.
    - Expensive.
  - Bottom up approach:
    - Leaf nodes are uniquely defined by their childmasks (8-bit mask where bit  $i$  tells us if child  $i$  contains geometry). Then at most  $2^8 = 256$  unique leaf nodes in an SVO.
    - First step: merge identical leaves.
    - Proceed to level above, identifying nodes that have identical childmasks and identical pointers (roots of identical subtrees that can be merged).
    - Iteratively perform previous step.

# High resolution sparse voxel DAGs



**Figure 1:** *The EPICCITADEL scene voxelized to a  $128K^3$  ( $131\,072^3$ ) resolution and stored as a Sparse Voxel DAG. Total voxel count is 19 billion, which requires 945MB of GPU memory. A sparse voxel octree would require 5.1GB without counting pointers. Primary shading is from triangle rasterization, while ambient occlusion and shadows are raytraced in the sparse voxel DAG at 170 MRays/sec and 240 MRays/sec respectively, on an NVIDIA GTX680.*

Kampe et al., 2013