

Data structures

CS594: Big Data Visualization & Analytics

Fabio Miranda


<https://fmiranda.me>

Overview

- Spatial data structures:
 - Uniform grid
 - Nested grids
 - Quadtree / octree
 - K-d tree
 - BVH
 - ...
- Visualization data structures:
 - Immens
 - Nanocubes
 - TopKube
 - ...

How to efficiently organize objects?

- 2D/3D data contains spatial information.
- How to perform queries when there are thousands / millions of objects (points, polygons)?
 - Ray-scene intersection.
 - Proximity queries
 - Point in polygon.
 - Range query.




Ray-scene intersection

- Given a scene with n primitives and a ray r , find the closest point of intersection of r with the scene.

```
function intersectObjects(ray, scene) {  
  for(var i=0; i < scene.objects.length; i++) {  
    var object = scene.objects[i];  
    var dist = intersection(ray, object);  
    // ...  
  }  
}
```

- Complexity: $O(n)$
- How to do better?

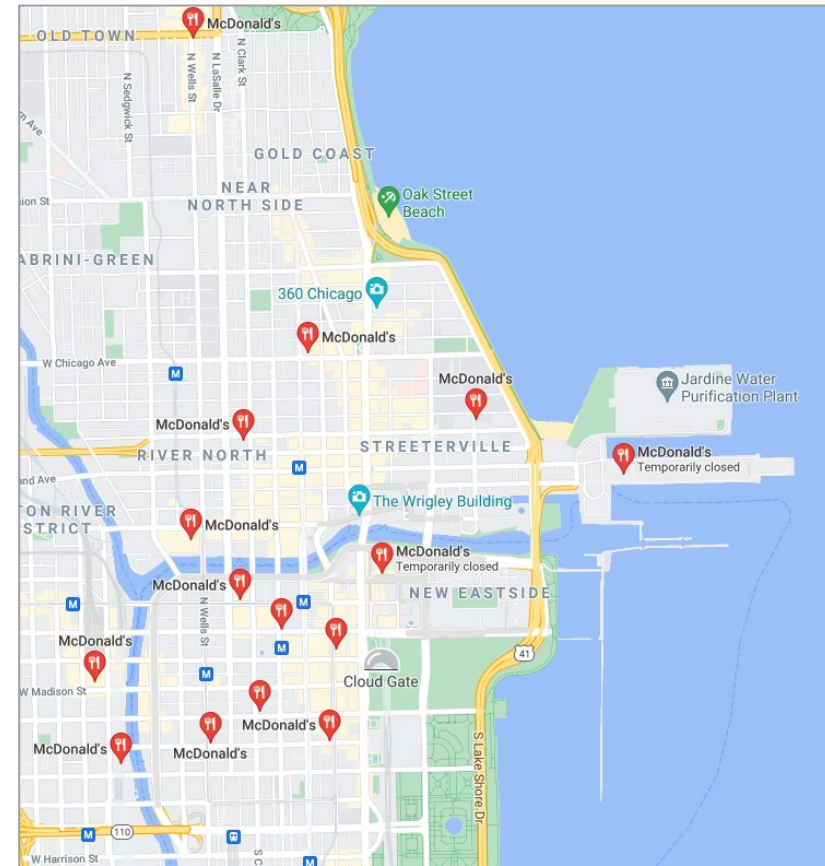


Proximity query

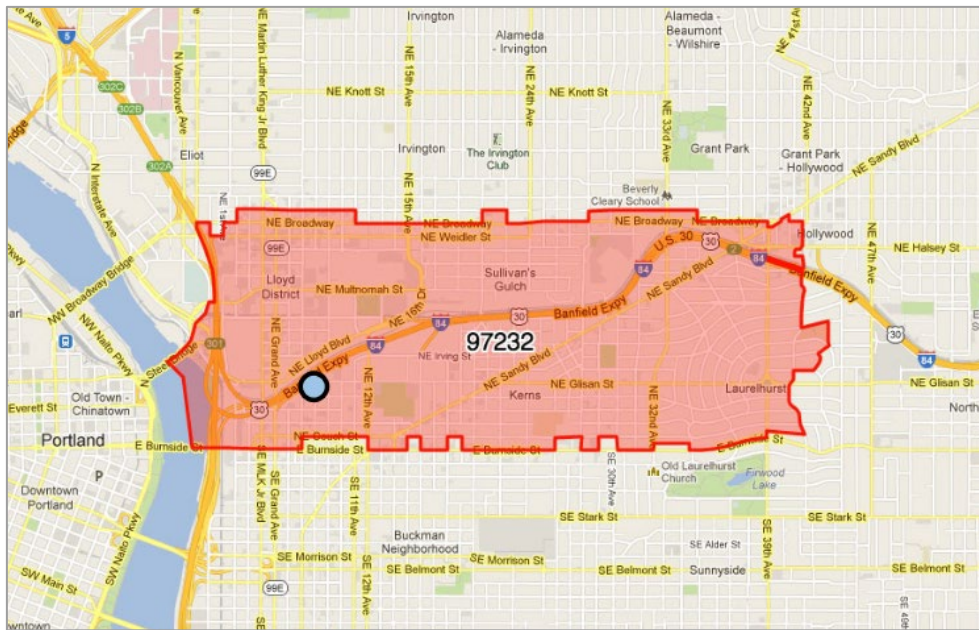
- Query based on proximity.
- “What is the closest McDonald’s?”

```
function findPlaces(query, scene) {  
  for(var i=0; i < scene.places.length; i++) {  
    var place = scene.places[i];  
    var dist = satisfyQuery(query, place);  
    // ...  
  }  
}
```

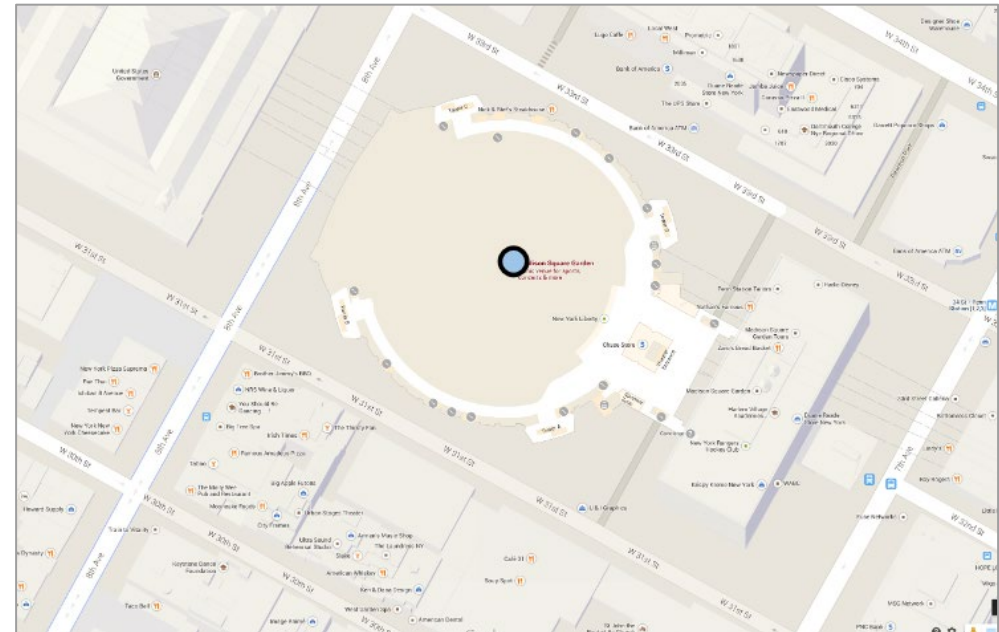
- Complexity: $O(n)$
- How to do better?



Point in polygon



What is the zip code for this complaint?



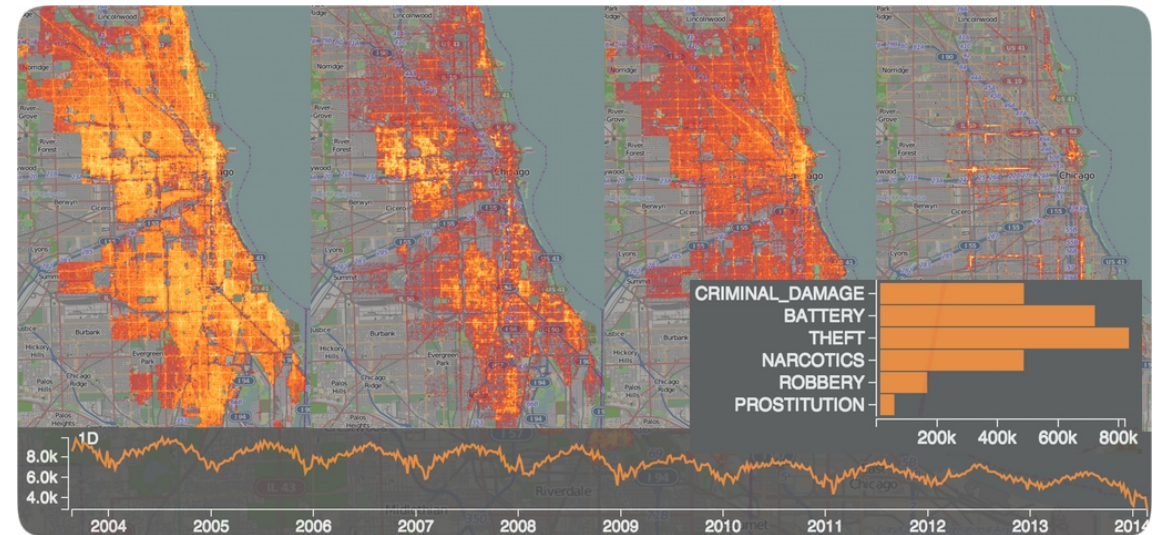
Am I inside a specific building?

Aggregation

- Aggregate spatiotemporal points.
- “Number of tweets in a region?”

```
function aggregate(points) {  
  var map = AggregatedData(width, height, 0);  
  for (var i=0; i < points.length; i++) {  
    var coord = getCoord(points[i]);  
    map.add(coord, 1);  
  }  
}
```

- Complexity: $O(n)$
- How to do better?



Time complexity

- Ray-scene intersection: $O(n)$
- Proximity query: $O(n)$
- Point in polygon: $O(n)$
- Aggregation: $O(n)$

How to reduce the time complexity?

Motivation

- Expensive operations (ray tracing, query).
 - Complex datasets (millions of objects).
 - Large number of operations (hundreds of millions per second).
- Reduce complexity through pre-processing data.
 - Spatial data structures: structures of objects in space.
 - Eliminate candidates as early as possible.
 - Reduce complexity to $O(\log n)$ on average.
 - Worst case complexity still $O(n)$.

Spatial data structures

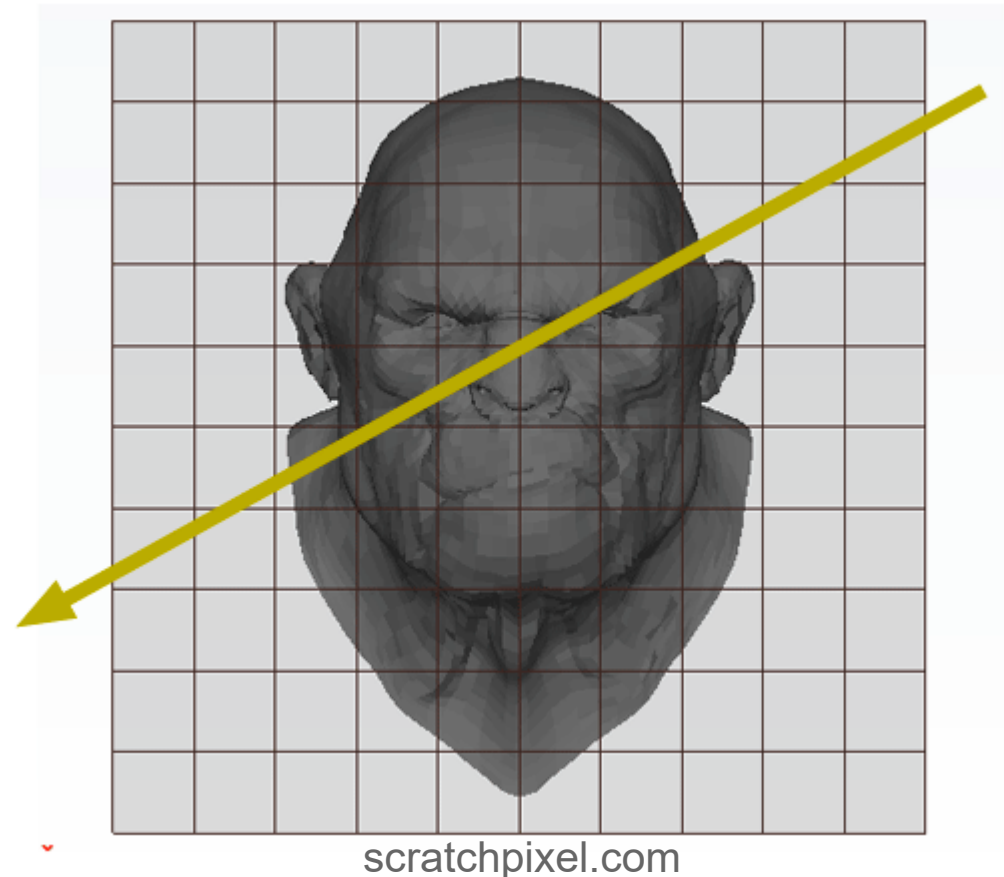
- Data structures to accelerate queries of the kind:
“I’m here, which object is around me?”
- Partition space or set of objects.
- Tasks:
 1. Construction / update:
 - Pre-processing for static parts of the scene.
 - Update for moving parts of the scene.
 2. Access:
 - Optimize so it is done as fast as possible.

Spatial data structures


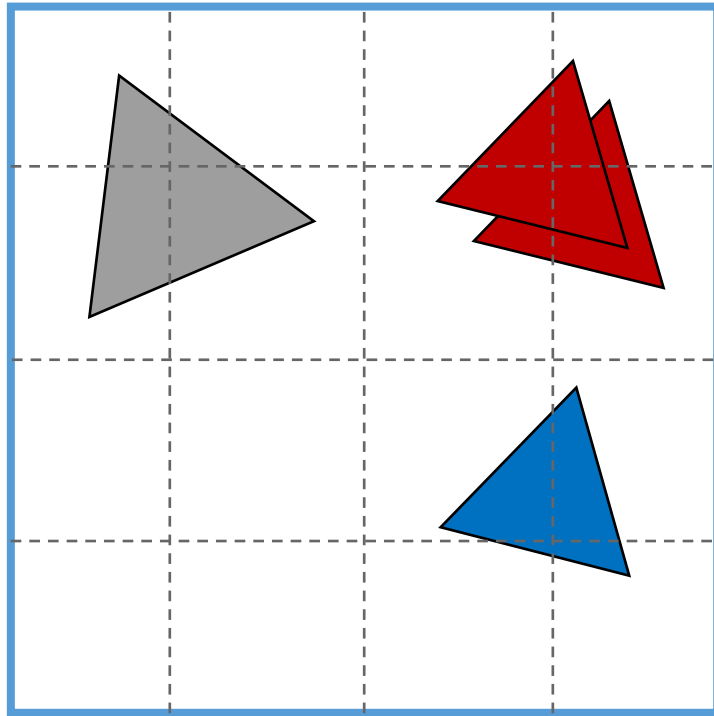
- Uniform grid: 2D/3D data, uniform distribution.
- Quadtree: 2D data, non-uniform distribution.
- Octree: 3D data, non-uniform distribution.
- KD-tree: 2D/3D data, avoid empty cells.

Uniform grid

- Partition space into equal-sized volumes (i.e., voxels).
- Each cell will contain objects that overlap the voxel.
- Good for uniform data (points are evenly distributed in space).
- Fast construction and queries.



Uniform grid



Uniform grid: construction and query

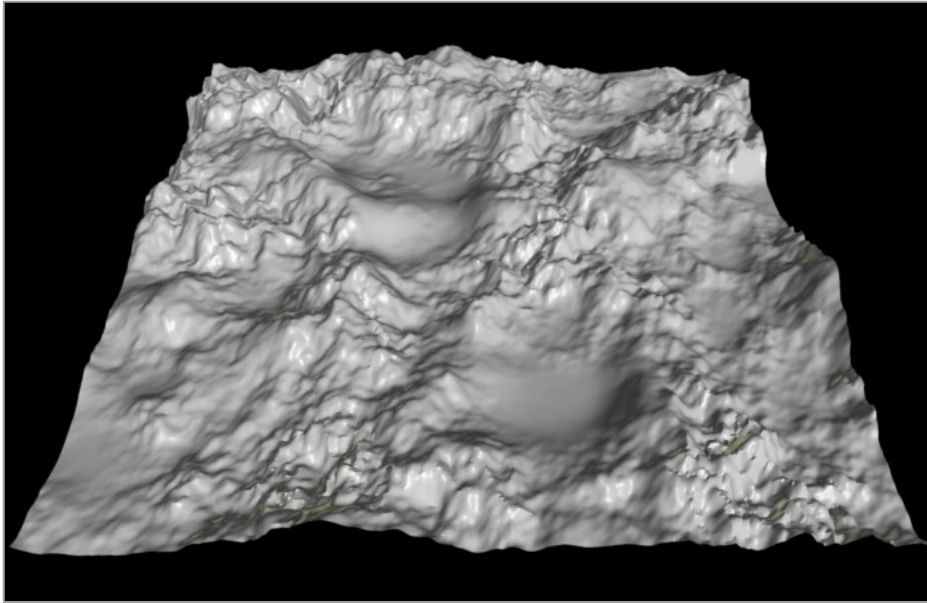
- Array of 3D voxels
 - Each voxel: list of pointers to colliding objects.
- Indexing function:
 - 3D point \rightarrow cell index (constant time!)
- Construction:
 - Initialize cells for grid with size $w * h$
 - For each object $p(x, y)$:
 - Compute grid cell using (x, y) .
 - Store p in cell.
- Query:
 - For query rectangle $(x_1, y_1) \times (x_2, y_2)$:
 - Compute subgrid for (x_1, y_1) and (x_2, y_2) .
 - For all cells inside subgrid, report all objects.
 - For all cells on the border of the subgrid, test objects against rectangle.

Uniform grid: complexity

- Build time: $O(n)$
- Space: $O(w * h) + O(n)$
- Query: $O(k)$

Uniform grid: complexity

- When uniform grids work well? Uniform distribution of objects.




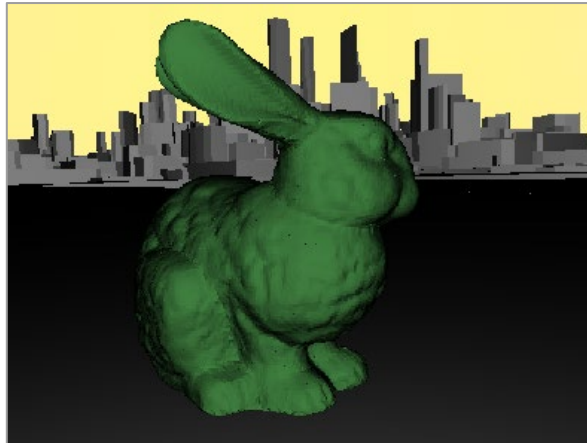

Mitsuba renderer



peterguthrie.net

Uniform grid: drawbacks

- When uniform grids do not perform well? Non-uniform distribution of objects.
- “Teapot in a stadium” problem: uniform grids cannot adapt to local density of objects.

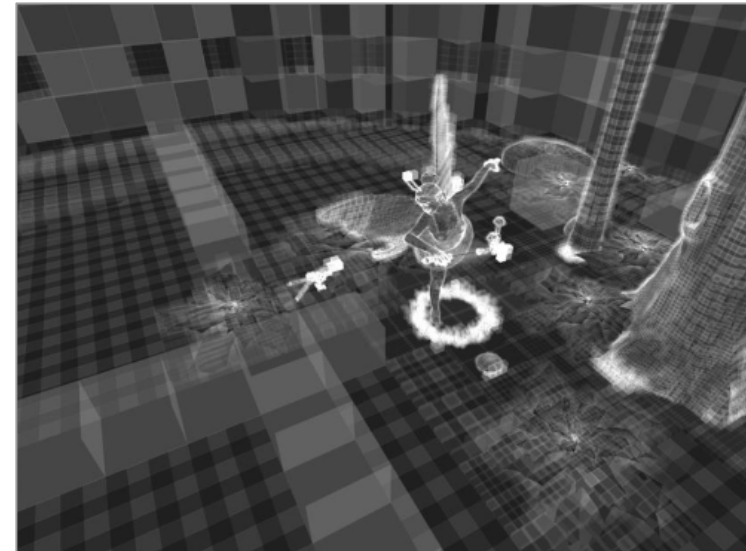
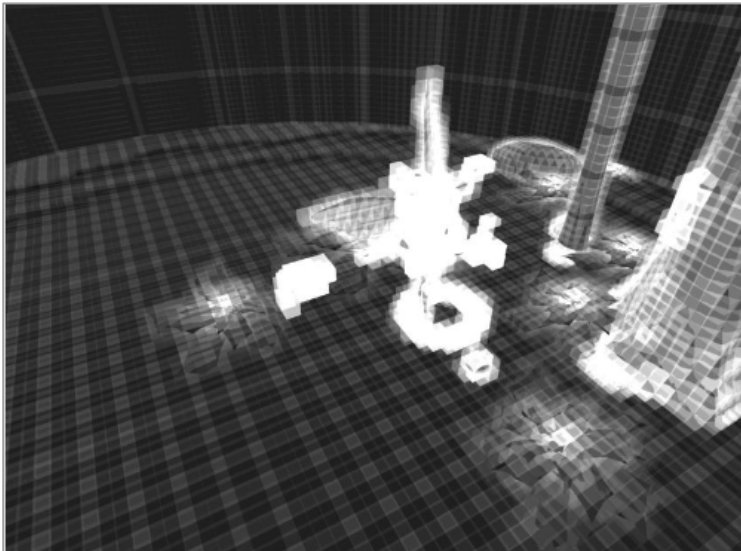


Uniform grid: drawbacks

- Assumes objects uniformly distributed in space.
- What happens when assumption does not hold?
 - Many empty cells.
 - Few cells with too many points.
- Change cell size?
 - Too small: memory occupancy too large.
 - Too big: too many objects in one cell.

Nested grids



- Possible solution to “teapot in a stadium” problem.
- Hierarchy of uniform grids: each cell is itself a grid.
- Fast building & traversal.



Philipp Slusallek

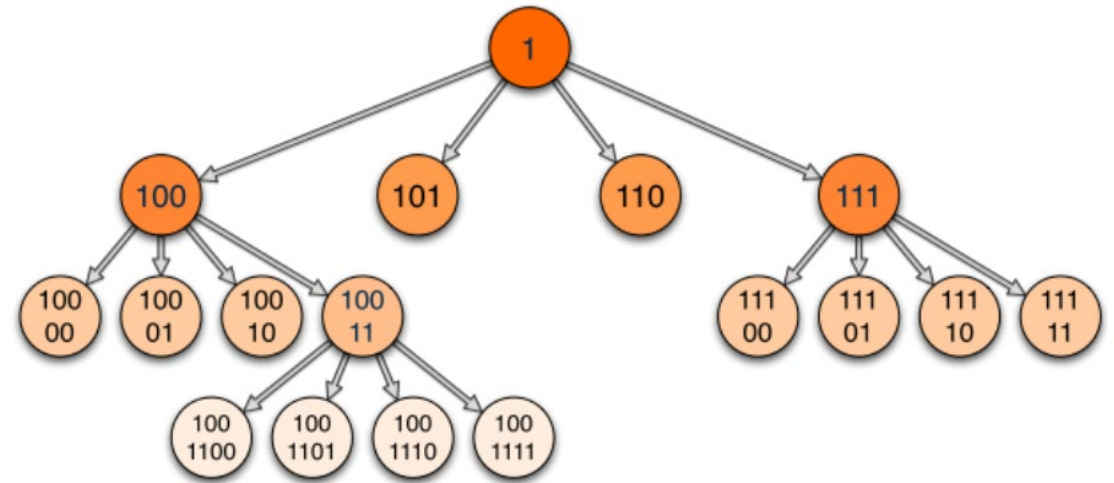
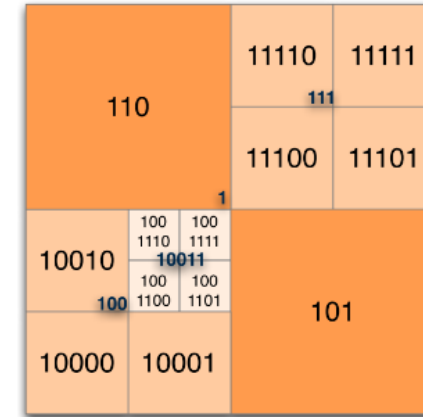
Quadtree

- Hierarchical structure that stores regular grids at each level.
- Adaptive subdivision: adjust depth to local scene complexity.



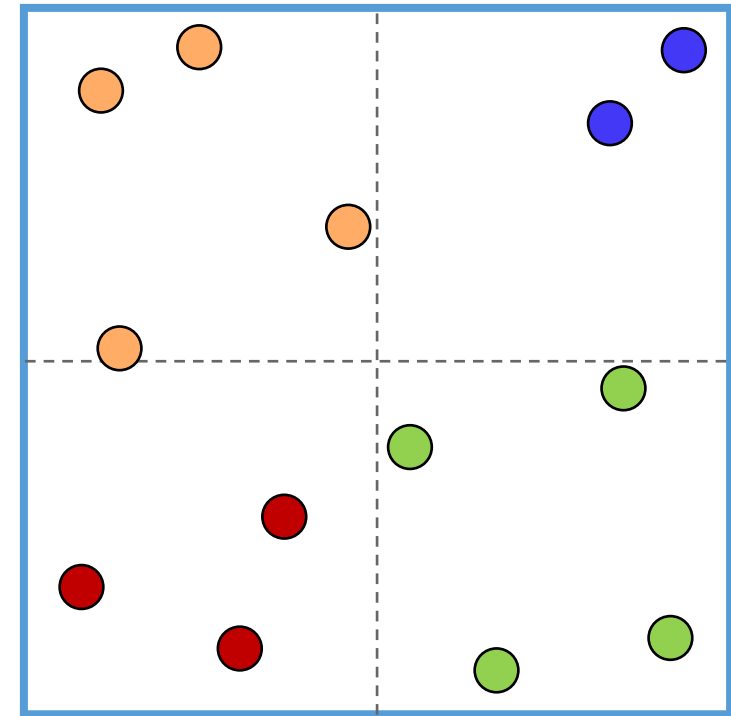
Quadtree

- Rooted tree in which every internal node has four children.
- Every node corresponds to a square.
- Tree: branching factor 4 or 8.
- Each node: splits into all dimensions at once (in the middle).
- Construction: continue splitting until end nodes have few objects (or limit level reached).




Quadtree: construction

- Split the top level.




Quadtree: construction

- Split the top level.
- Can we stop?




Quadtree: construction

- Split the top level.
- Can we stop? No, split the next level.




Quadtree: construction

- Split the top level.
- Can we stop? No, split the next level.
- Split top-left.



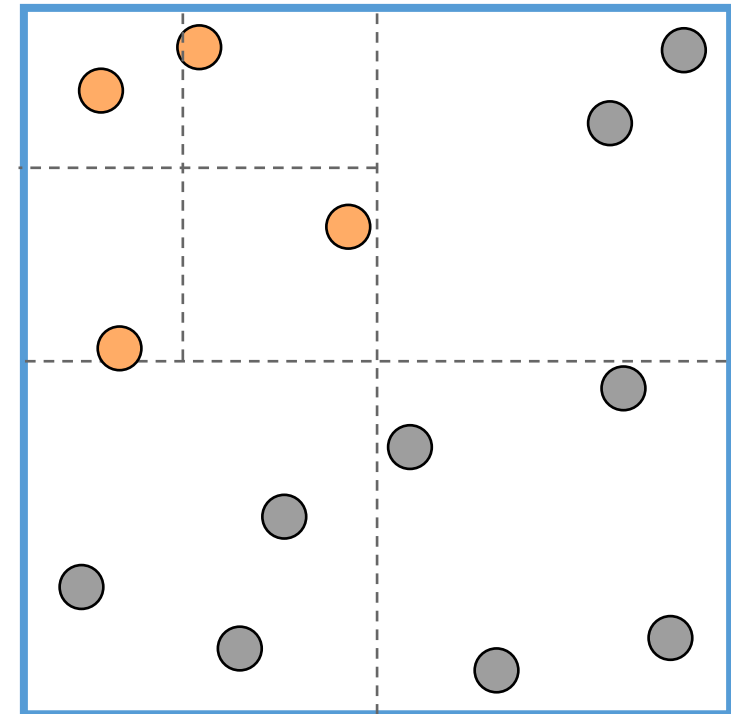
Quadtree: construction

- Split the top level.
- Can we stop? No, split the next level.
- Split top-left.
- Can we stop top-left?



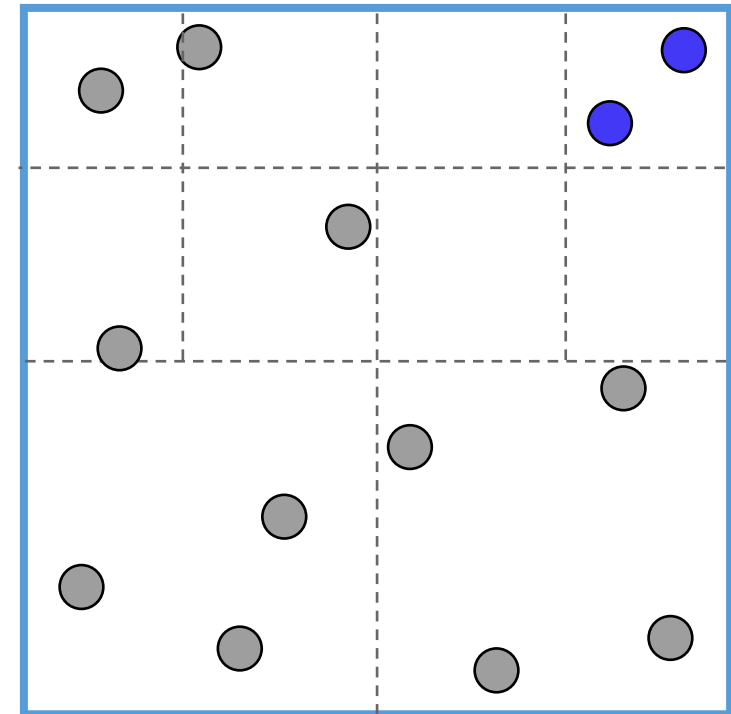
Quadtree: construction

- Split the top level.
- Can we stop? No, split the next level.
- Split top-left.
- Can we stop top-left? Yes.



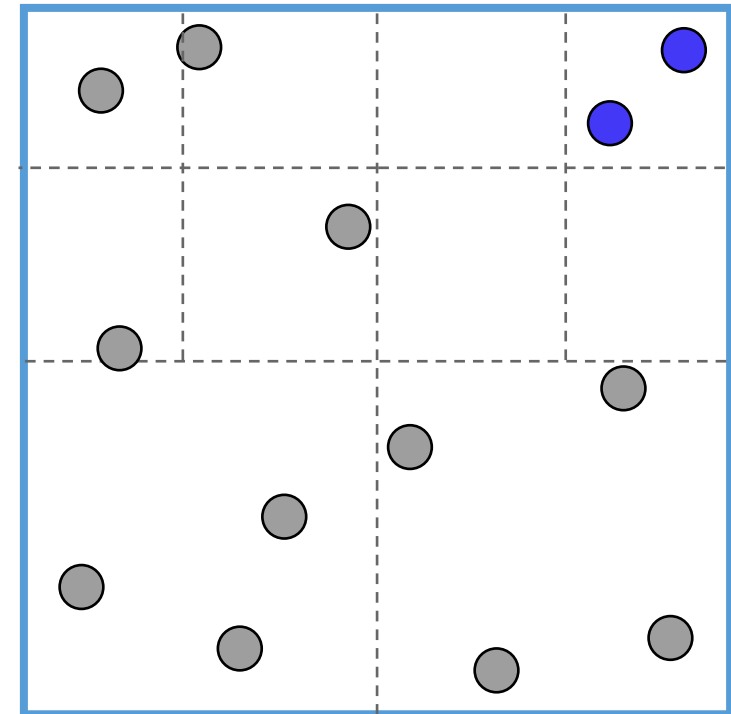
Quadtree: construction

- Split the top level.
- Can we stop? No, split the next level.
- Split top-left.
- Can we stop top-left? Yes.
- Split top-right.




Quadtree: construction

- Split the top level.
- Can we stop? No, split the next level.
- Split top-left.
- Can we stop top-left? Yes.
- Split top-right.
- Can we stop top-right?



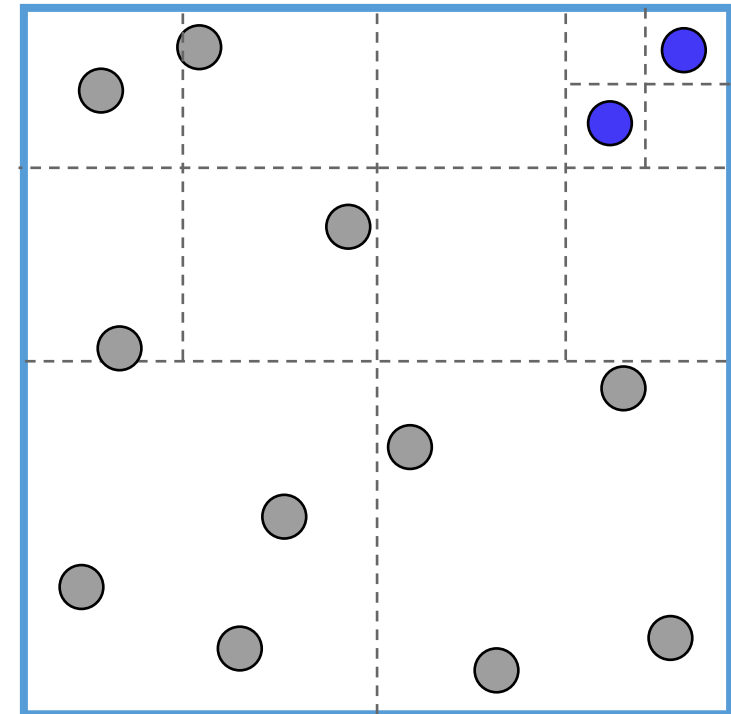
Quadtree: construction

- Split the top level.
- Can we stop? No, split the next level.
- Split top-left.
- Can we stop top-left? Yes.
- Split top-right.
- Can we stop top-right? No.



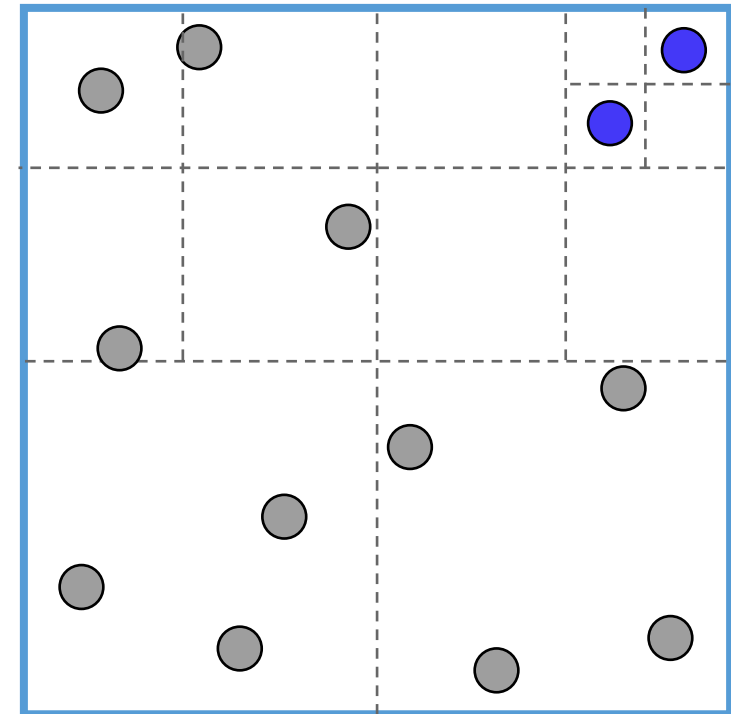
Quadtree: construction

- Split the top level.
- Can we stop? No, split the next level.
- Split top-left.
- Can we stop top-left? Yes.
- Split top-right.
- Can we stop top-right? No.
- Split top-right.



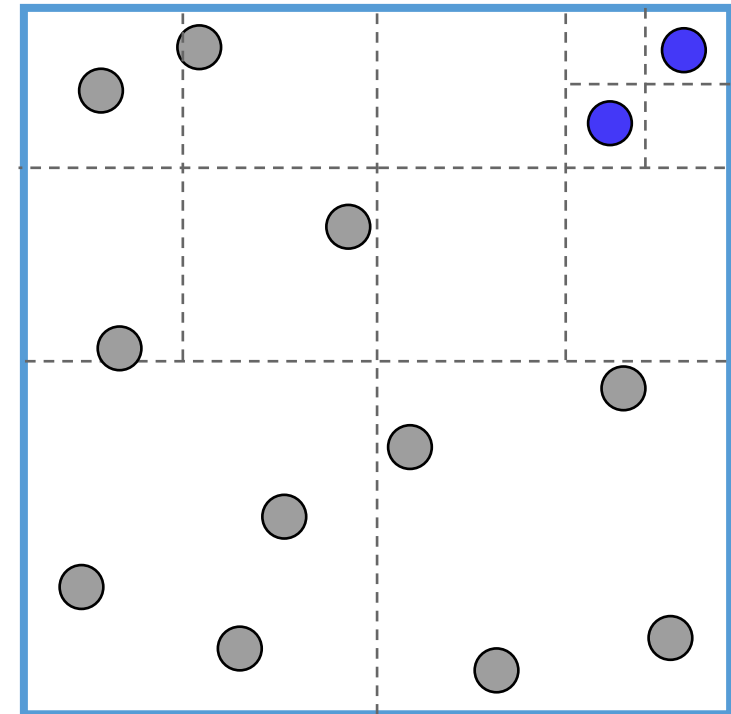
Quadtree: construction

- Split the top level.
- Can we stop? No, split the next level.
- Split top-left.
- Can we stop top-left? Yes.
- Split top-right.
- Can we stop top-right? No.
- Split top-right.
- Can we stop top-right?



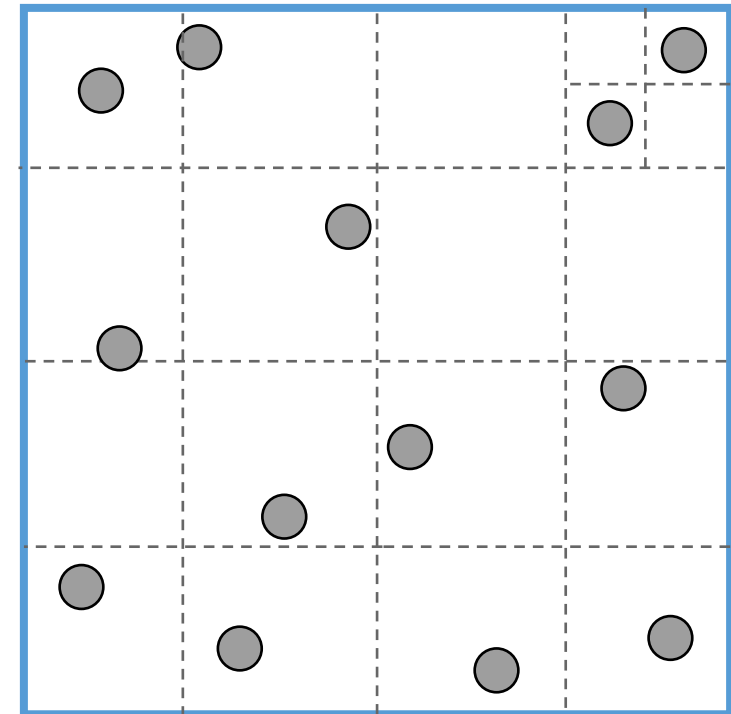
Quadtree: construction

- Split the top level.
- Can we stop? No, split the next level.
- Split top-left.
- Can we stop top-left? Yes.
- Split top-right.
- Can we stop top-right? No.
- Split top-right.
- Can we stop top-right? Yes.

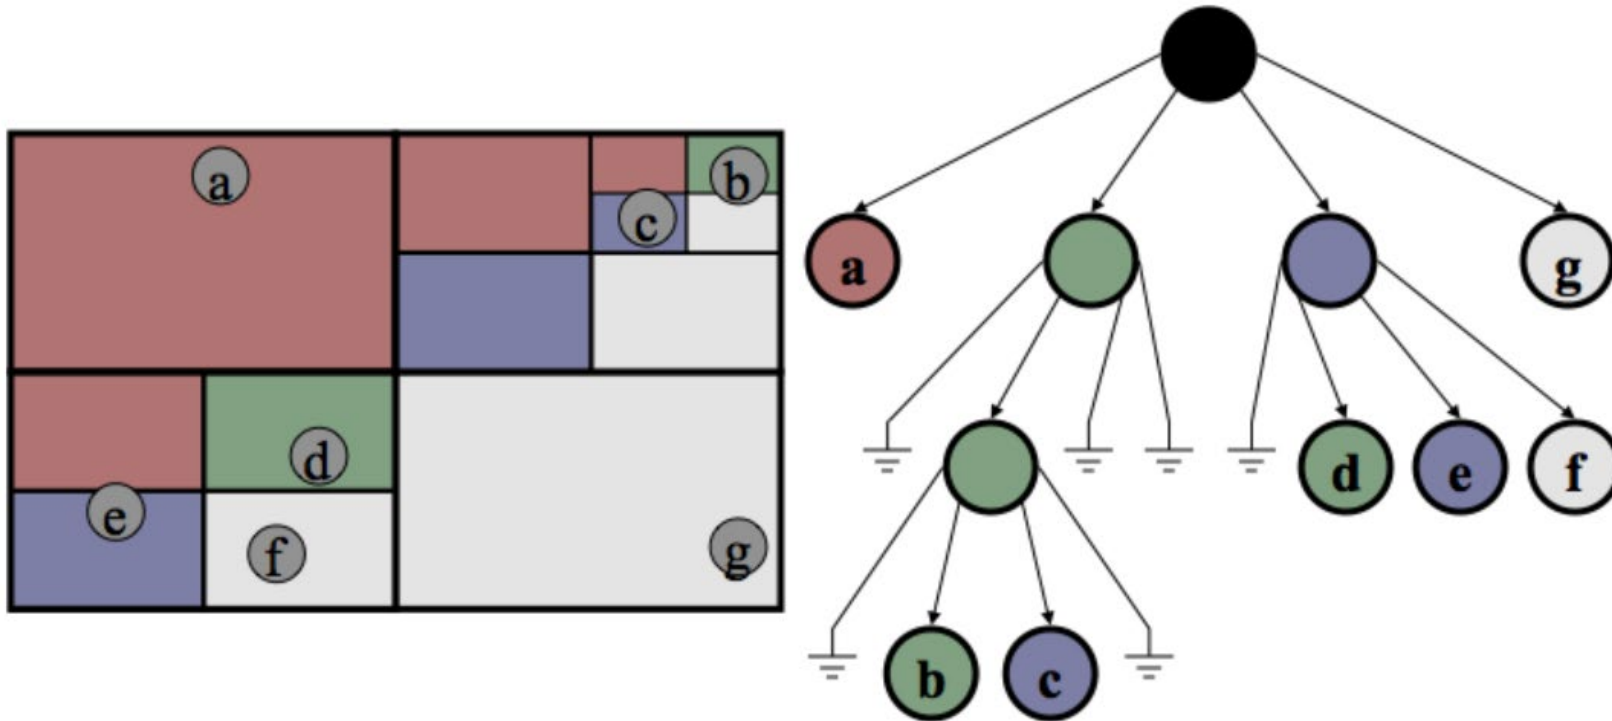


Quadtree: construction

- Split the top level.
- Can we stop? No, split the next level.
- Split top-left.
- Can we stop top-left? Yes.
- Split top-right.
- Can we stop top-right? No.
- Split top-right.
- Can we stop top-right? Yes.



Quadtree: construction



Quadtree: construction

- Construction:
 - Input: set of objects P inside a square $S (x_1, y_1) \times (x_2, y_2)$, tree node v
 - If $|P| \leq 1$:
 - Quadtree consists of a single leaf with P .
 - Else:
 - P_{00} : set of points that fall in the bottom-left corner of S .
 - P_{01} : set of points that fall in the bottom-right corner of S .
 - ...
 - v_{00} : node with points of P_{00} .
 - v_{01} : node with points of P_{01} .
 - ...
 - Append $v_{00}, v_{01}, v_{10}, v_{11}$ to v .

Quadtree: query

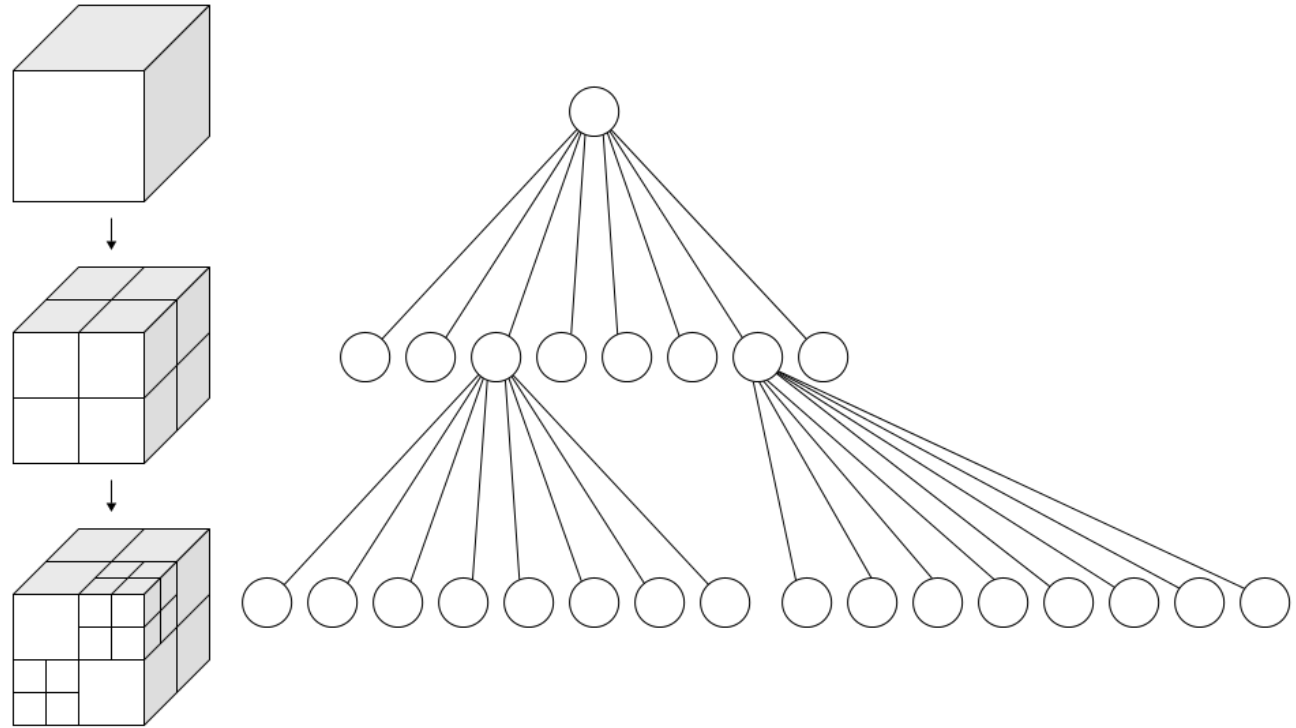
- Query:
 - Input: range query r $(x_1, y_1) \times (x_2, y_2)$, tree node v .
 - If v is a leaf:
 - Search points of v inside range r .
 - If v_{00} inside range r :
 - Query(v_{00} , r)
 - If v_{01} inside range r :
 - Query(v_{01} , r)
 - If v_{10} inside range r :
 - Query(v_{10} , r)
 - If v_{11} inside range r :
 - Query(v_{11} , r)

Quadtree: complexity

- Build time: $O(n)$
- Space: $O(n)$
- Range query: $O(\sqrt{n} + k)$
- Leaf traversal: $O(\log n)$

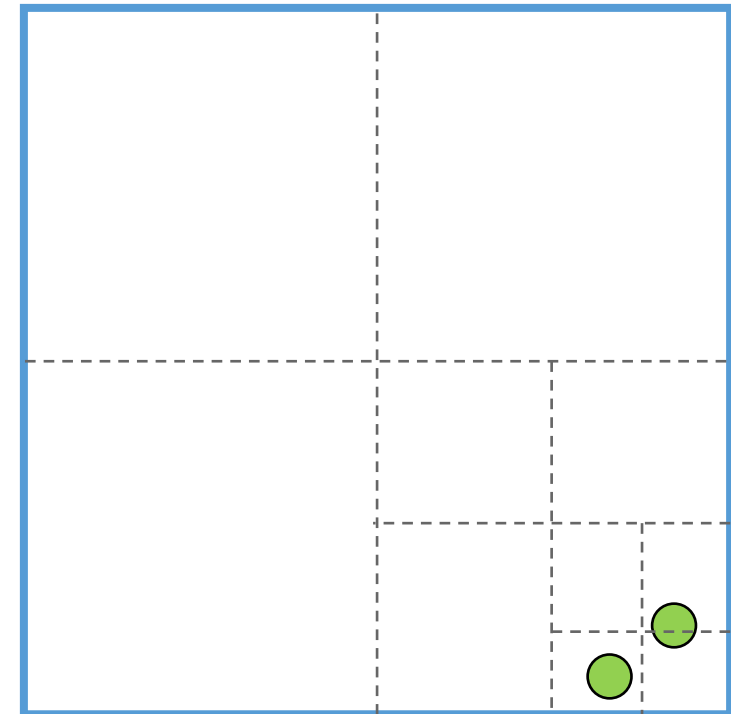
Octree

- Each inner node contains 8 equally sized voxels.
- A 3D quadtree.



Quadtree and octree: drawbacks

- Greater ability to adapt to location of scene geometry than uniform grid.
- But very long tree to store points that are concentrated in a small region.
- Many nodes will contain zero objects.



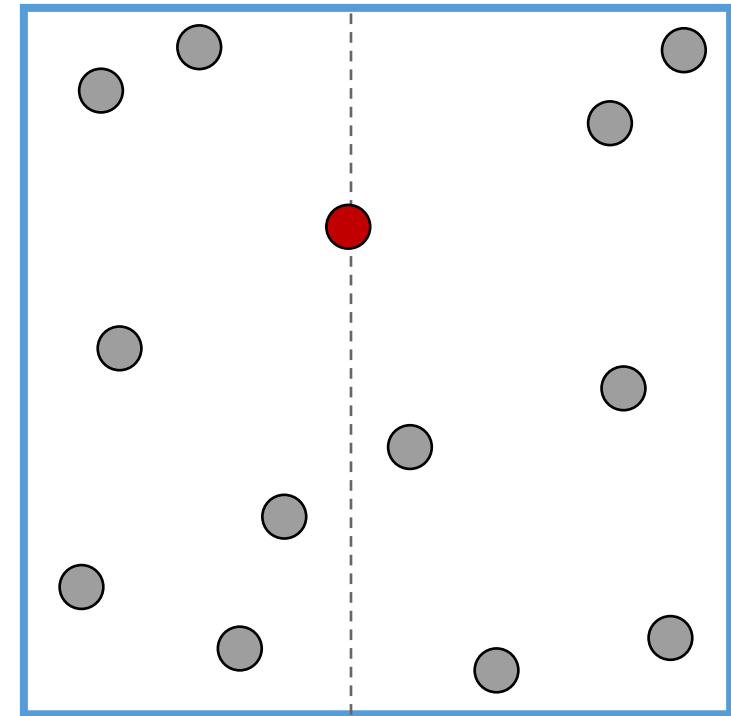
K-d tree



- Differently from quadtrees and octrees, k-d trees only split **one** dimension at each level.
- Where to split? Middle? Median? Proportional to surface area?
- At each level:
 - Quadtree creates 4 equal sized cells.
 - Octree creates 8 equal sized cells.
 - K-d tree creates 2 non-equal sized cells (2D case).

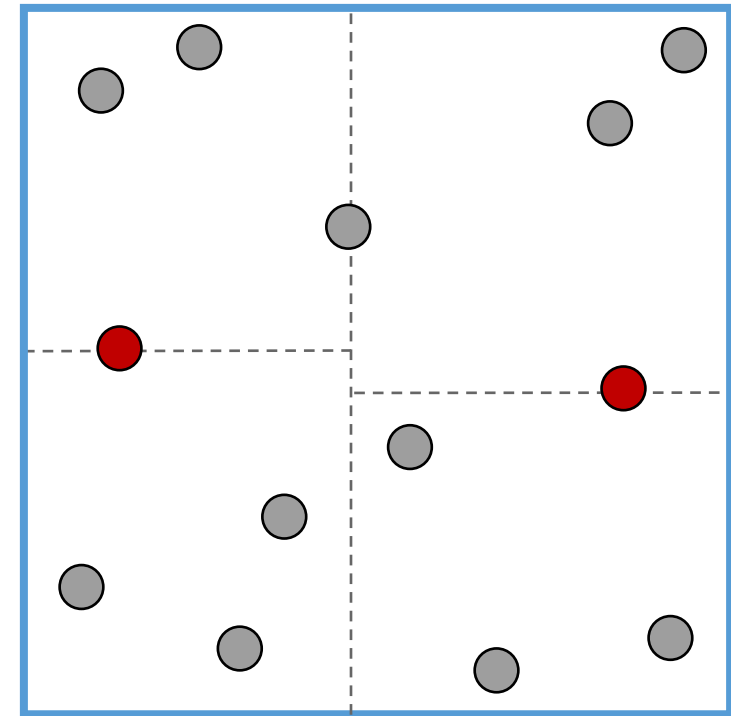
K-d tree: construction

- First split: x dimension (median point).



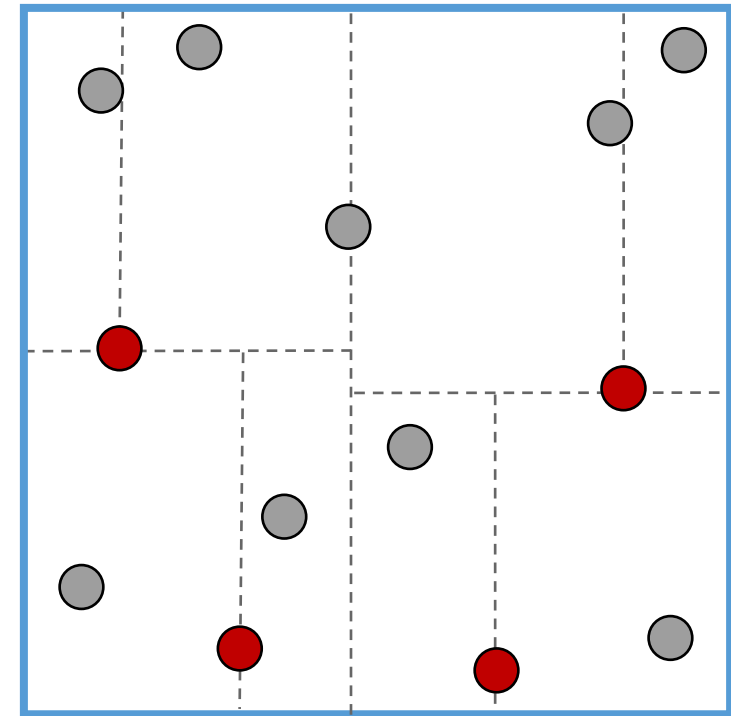
K-d tree: construction

- First split: x dimension (median point).
- Second split: y dimension.



K-d tree: construction

- First split: x dimension (median point).
- Second split: y dimension.
- Repeat, alternating split dimensions



K-d tree: construction

- Construction:
 - Input: set of objects P inside a square $S (x_1, y_1) \times (x_2, y_2)$, tree node v
 - If $|P| \leq 1$:
 - K-d tree consists of a single leaf with P .
 - Else:
 - If depth is even:
 - Split P into P_0 and P_1 , along a vertical line through the y axis.
 - Else:
 - Split P into P_0 and P_1 , along a vertical line through the x axis.
 - v_0 : $build(v, P_0, depth + 1)$.
 - v_1 : $build(v, P_1, depth + 1)$.
 - ...
 - Append v_0, v_1 to v .

K-d tree: query

- Query:
 - Input: range query $r (x_1, y_1) \times (x_2, y_2)$, tree node v .
 - If v is a leaf:
 - Search points of v inside range r .
 - If v_0 inside range r :
 - Query(v_0, r)
 - If v_1 inside range r :
 - Query(v_1, r)

K-d tree: complexity

- Build time: $O(n \log n)$
- Space: $O(n)$
- Range query: $O(\sqrt{n} + k)$
- Leaf traversal: $O(\log n)$

K-d tree and Scikit-learn

- K-nearest neighbors and neighbors within a radius:

```
from sklearn.neighbors import KDTree
import numpy as np
```

```
rng = np.random.RandomState(0)
X = rng.random_sample((1000, 2))
tree = KDTree(X, leaf_size=2)
dist, ind = tree.query(X[:1], k=3)
```

```
from sklearn.neighbors import KDTree
import numpy as np
```

```
rng = np.random.RandomState(0)
X = rng.random_sample((1000, 2))
tree = KDTree(X, leaf_size=2)
points = tree.query_radius(X[:1], r=0.3)
```

- Kernel density estimation:

```
from sklearn.neighbors import KDTree
import numpy as np
```

```
rng = np.random.RandomState(0)
X = rng.random_sample((1000, 2))
tree = KDTree(X, leaf_size=2)
estimate = tree.kernel_density(X[:3], h=0.1, kernel='gaussian')
```


K-d tree and Scikit-learn

- tSNE


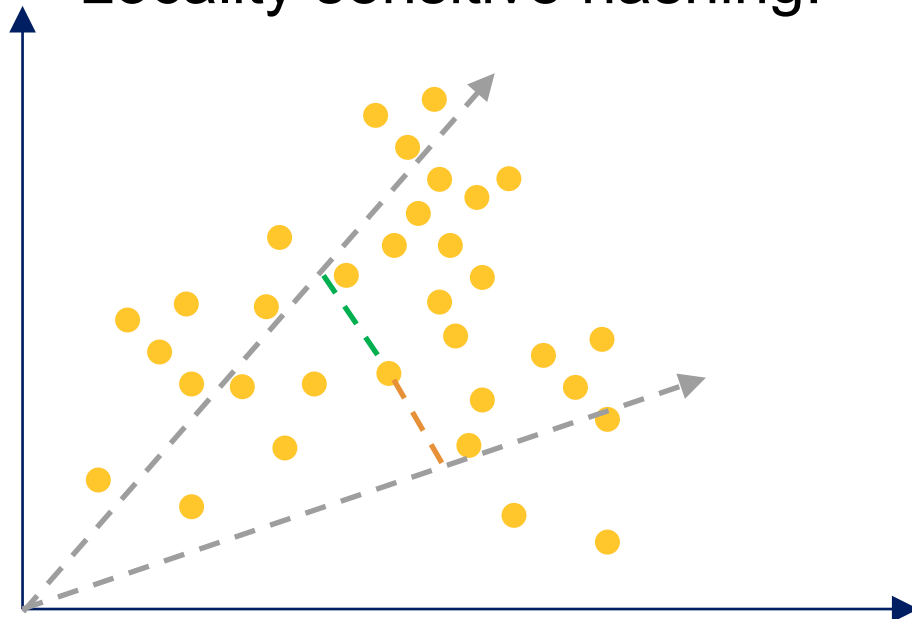
```
import numpy as np
from sklearn.neighbors import KNeighborsTransformer
from sklearn.pipeline import make_pipeline

rng = np.random.RandomState(0)
X = rng.random_sample((1000, 2))
transformer = make_pipeline(
    KNeighborsTransformer(n_neighbors=n_neighbors, mode='distance', metric=metric),
    TSNE(metric='precomputed', **tsne_params)
)

transformer.fit_transform(X)
```

Annoy and Scikit-learn

- Annoy: approximate nearest neighbors C++ library with Python bindings.
- Locality sensitive hashing:



$$\alpha_{1,2} = \cos^{-1}\left(\frac{\vec{v}_1 \cdot \vec{v}_2}{|\vec{v}_1||\vec{v}_2|}\right)$$

Annoy and Scikit-learn

- tSNE

```
class AnnoyTransformer(TransformerMixin, BaseEstimator):
    def fit(self, X):
        self.n_samples_fit_ = X.shape[0]
        self.annoy_ = annoy.AnnoyIndex(X.shape[1], metric=self.metric)
        for i, x in enumerate(X):
            self.annoy_.add_item(i, x.tolist())
        self.annoy_.build(self.n_trees)
        return self
    (...)
rng = np.random.RandomState(0)
X = rng.random_sample((1000, 2))
transformer = make_pipeline(
    AnnoyTransformer(n_neighbors=n_neighbors, metric=metric),
    TSNE(metric='precomputed', **tsne_params)
)

embedded = transformer.fit_transform(X)
```

Annoy and Scikit-learn

- tSNE

```
class AnnoyTransformer(TransformerMixin, BaseEstimator):
    def fit(self, X):
        self.n_samples_fit_ = X.shape[0]
        self.annoy_ = annoy.AnnoyIndex(X.shape[1], metric=self.metric)
        for i, x in enumerate(X):
            self.annoy_.add_item(i, x.tolist())
        self.annoy_.build(self.n_trees)
        return self
    (...)
rng = np.random.RandomState(0)
X = rng.random_sample((1000, 2))
transformer = make_pipeline(
    AnnoyTransformer(n_neighbors=n_neighbors, metric=metric),
    TSNE(metric='precomputed', **tsne_params)
)

embedded = transformer.fit_transform(X)
```

TSNE with AnnoyTransformer: 30.225 sec
TSNE with KNeighborsTransformer: 64.845 sec

Summary

- Choose the right structure considering the operations and data.
- Uniform grid:
 - The most parallelizable (to update, construct, use).
 - Constant time access (best!).
 - Quadratic / cubic space (2D, 3D).
 - Good performance under uniform distribution of objects.
- Quadtree, octree, k-d tree:
 - Compact.
 - Simple.
 - Non-constant accessing time.
 - Good performance under non-uniform distribution of objects.

Data structures for visualization

- Immens
- Nanocube
- TopKube
- Learned cubes