

Spatial data processing

CS594: Big Data Visualization & Analytics

Fabio Miranda

<https://fmiranda.me>

Overview



- Spatial data
- Spatial queries
- Spatial indices

Spatial data



flickr



twitter



Infrastructure

Environment

Social media

Spatial data

- Spatial attributes
 - 2D: (x, y)
 - 3D: (x, y, z)
- Spatial data primitives:
 - Points
 - Lines
 - Polygons
- Other attributes:
 - Time
 - ...



Spatial queries

- Spatial selection
- Spatial joins
- Spatial aggregation

Spatial selection queries

Select spatial objects that satisfy a spatial constraint.

```
SELECT *  
FROM taxi T  
WHERE T.pickup inside Lower Manhattan
```



Spatial selection queries

Select spatial objects that satisfy a spatial constraint.

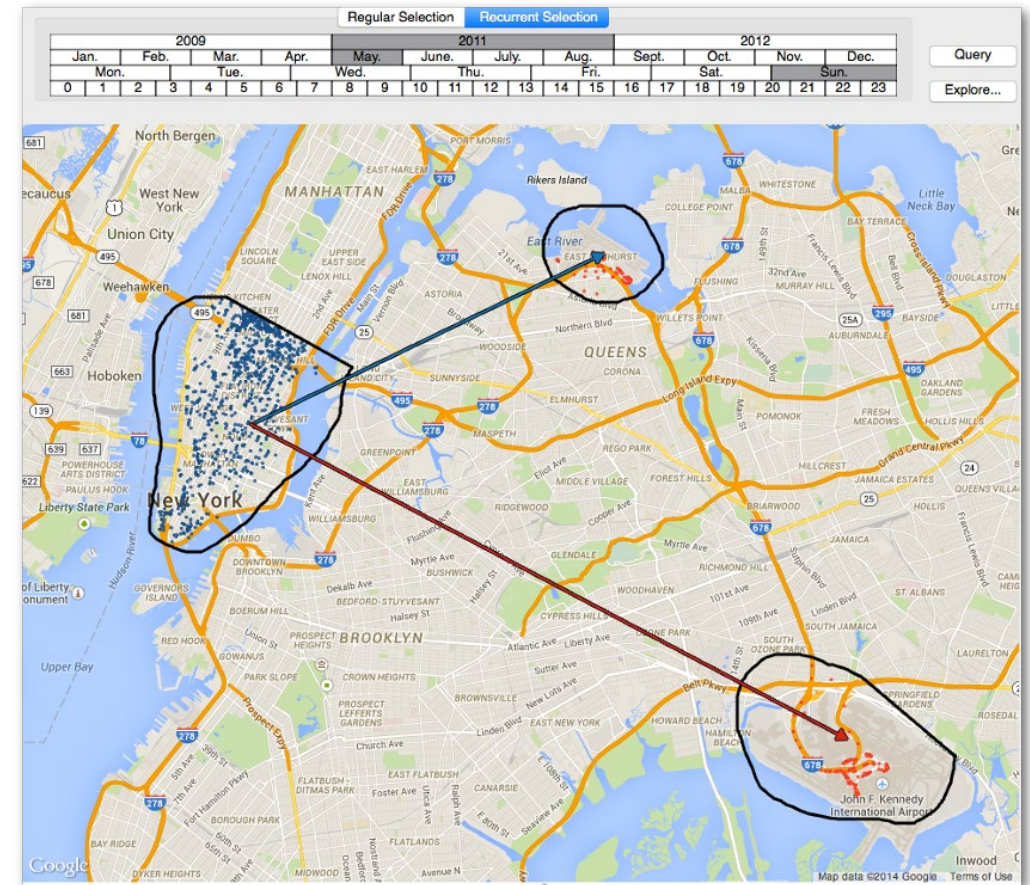
```
SELECT *  
FROM taxi T  
WHERE T.pickup inside Lower Manhattan  
AND (T.dropoff inside JFK OR  
T.dropoff inside LGA)
```



Spatial selection queries

Select spatial objects that satisfy a spatial constraint and other constraints.


```
SELECT *  
FROM taxi T  
WHERE T.pickup inside Lower Manhattan  
AND (T.dropoff inside JFK OR  
T.dropoff inside LGA)  
AND T.picktime in May 2011
```



Spatial joins

Select pairs of spatial objects satisfying a spatial constraint.


```
SELECT *  
FROM taxi T, neighborhoods N  
WHERE T.pickup inside N.polygon
```



Spatial aggregation queries

Aggregate data points over spatial regions.


```
SELECT COUNT(*)  
FROM taxi T, neighborhoods N  
WHERE T.pickup inside N.polygon  
GROUP BY N.id
```



Spatial aggregation queries

Aggregate data points over spatial regions.

```
SELECT AVG(T.duration)
FROM taxi T, neighborhoods N
WHERE T.pickup inside N.polygon
GROUP BY N.id
```



Nearest neighbor queries

Find nearest points.


```
SELECT TOP(10)
FROM lots B, crime C
ORDER BY DISTANCE(B.geometry,
                  C.location)
```



Spatial index

- How to speedup these queries?
 - Spatial index!

1-dimensional data




1-D range search: Find points between x and y

Unordered list: Fast insert $O(1)$, slow range search $O(n)$


Ordered list: Slow insert $O(n)$, binary search for x and y to do range search $O(\log n)$.

1-dimensional data




Find point with value 5

1-dimensional data: binary search




Find point with value 5

Binary Search Tree




1-dimensional data: binary search




Find point with value 5

Binary Search Tree




1-dimensional data: binary search




Find point with value 5

Binary Search Tree




1-dimensional data: binary search




Find point with value 5

Binary Search Tree




1-dimensional data: binary search




Find point with value 5

Binary Search Tree



1-dimensional data: hash function



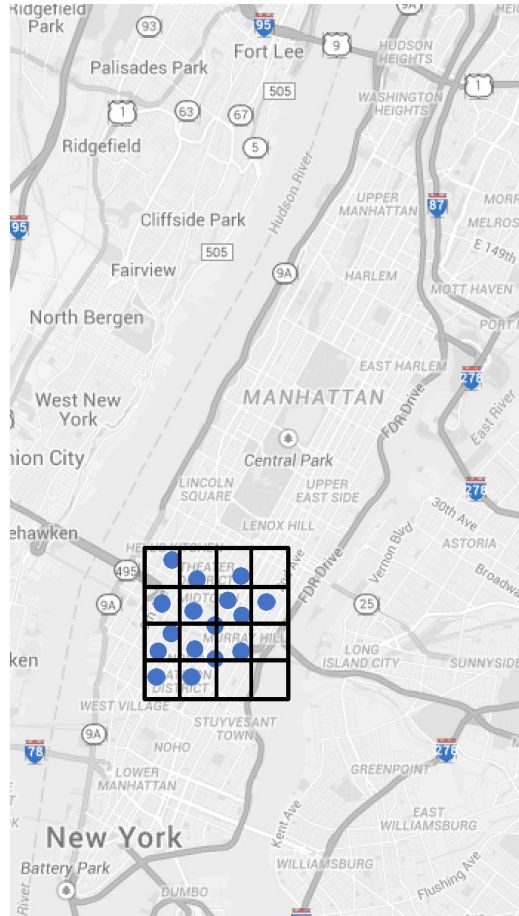
- Create bins using a hash function.
- Query:
 - Identify bin(s) satisfying query constraint.
 - Search within bin.

2-dimensional data




- Kd-Tree:
 - Extension of a binary search tree to higher dimensions.
 - Supports k-dimensional tree.
- Grid index:
 - Extension of hash index to higher dimensions.
 - Hash function is defined by a grid.
 - Overlay a grid covering the spatial region → assign objects to different grid cells.

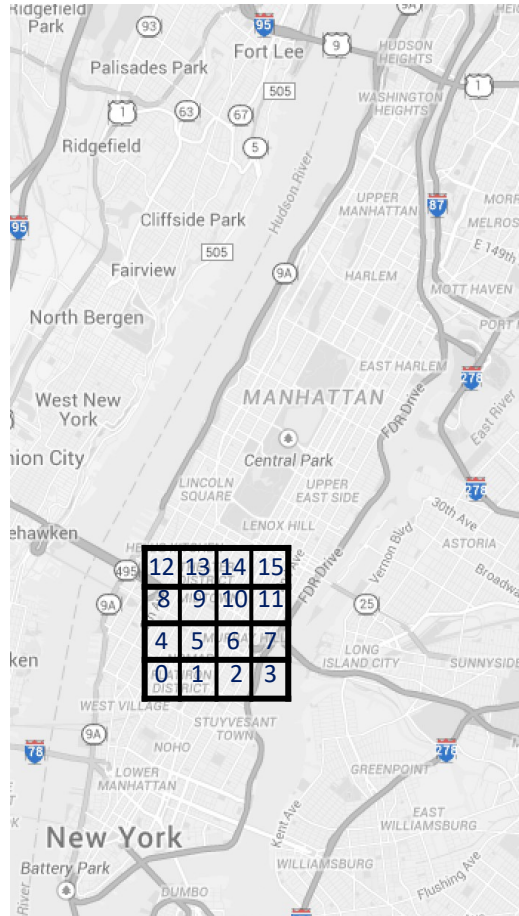
2-dimensional example



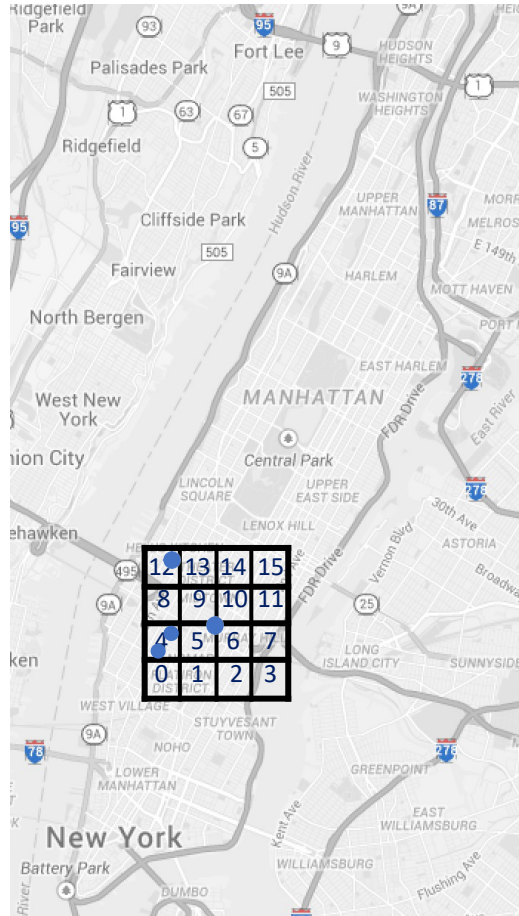
2-dimensional example



2-dimensional example

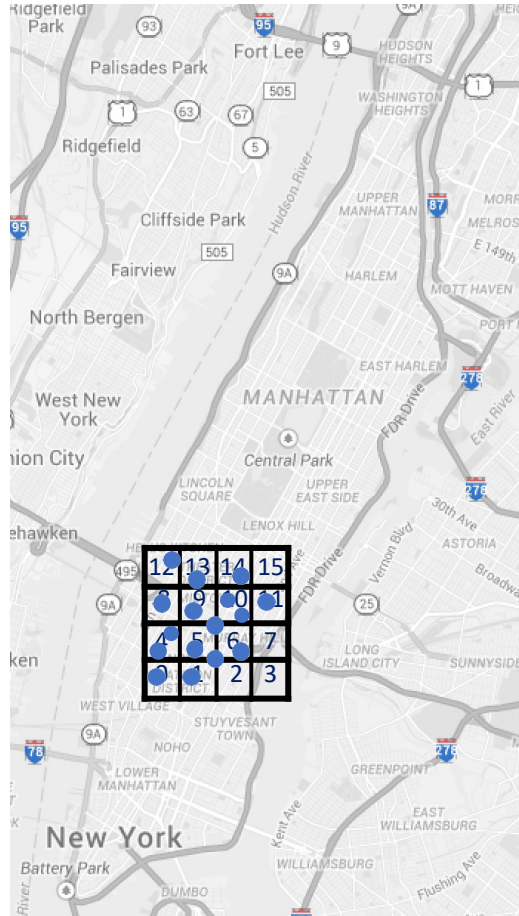


2-dimensional example



| | |
|----|-----|
| 0 | |
| 1 | |
| 2 | |
| 3 | |
| 4 | 3 4 |
| 5 | 10 |
| 6 | 10 |
| 7 | |
| 8 | |
| 9 | |
| 10 | |
| 11 | |
| 12 | 1 |
| 13 | |
| 14 | |
| 15 | |

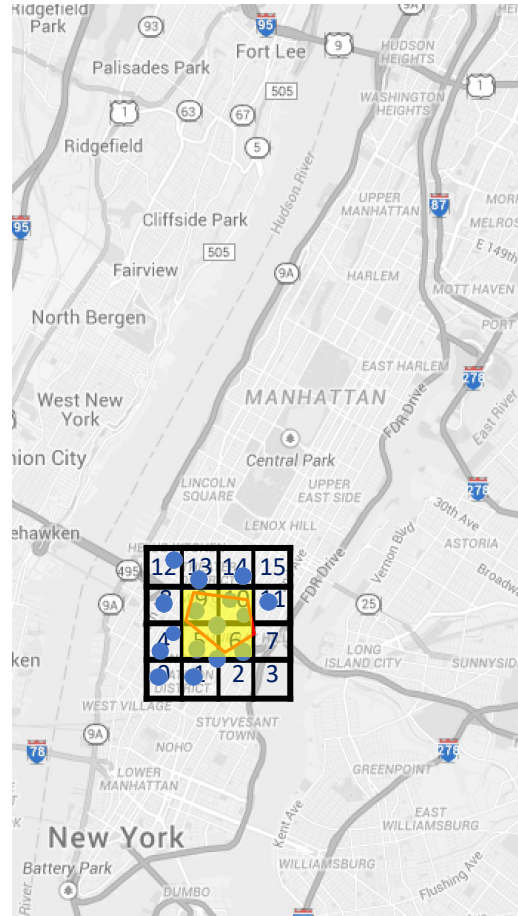
2-dimensional example



| | |
|----|----------|
| 0 | 5 |
| 1 | 9 11 |
| 2 | 11 |
| 3 | |
| 4 | 3 4 |
| 5 | 8 10 11 |
| 6 | 10 11 15 |
| 7 | |
| 8 | 2 |
| 9 | 7 |
| 10 | 12 14 |
| 11 | 16 |
| 12 | 1 |
| 13 | 6 |
| 14 | 14 |
| 15 | |

2-dimensional example

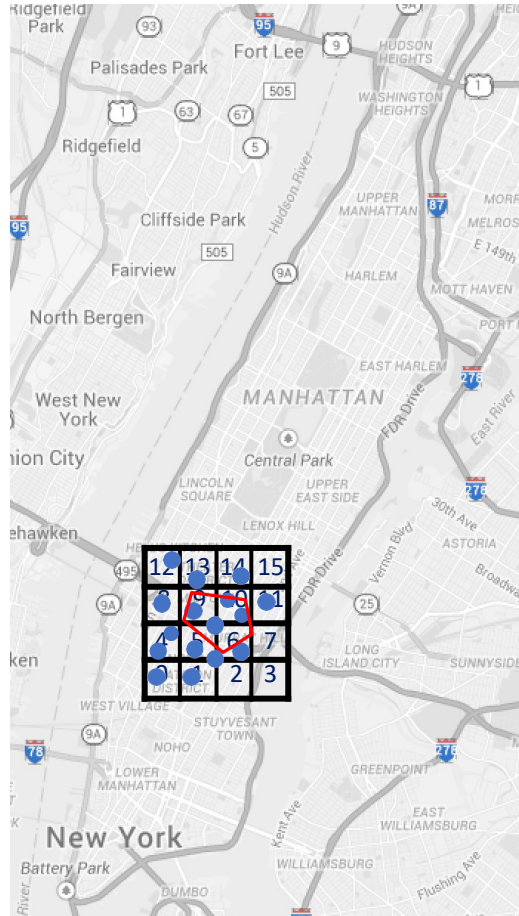
- Find Cells Intersected:
 - 5,6,9,10
- Test all points in these cells



| | |
|----|----------|
| 0 | 5 |
| 1 | 9 11 |
| 2 | 11 |
| 3 | |
| 4 | 3 4 |
| 5 | 8 10 11 |
| 6 | 10 11 15 |
| 7 | |
| 8 | 2 |
| 9 | 7 |
| 10 | 12 14 |
| 11 | 16 |
| 12 | 1 |
| 13 | 6 |
| 14 | 14 |
| 15 | |

2-dimensional example

- Find Cells Intersected:
 - 5,6,9,10
- Test all points in these cells




| | |
|----|----------|
| 0 | 5 |
| 1 | 9 11 |
| 2 | 11 |
| 3 | |
| 4 | 3 4 |
| 5 | 8 10 11 |
| 6 | 10 11 15 |
| 7 | |
| 8 | 2 |
| 9 | 7 |
| 10 | 12 14 |
| 11 | 16 |
| 12 | 1 |
| 13 | 6 |
| 14 | 14 |
| 15 | |

Spatial queries

- Spatial selection
- Spatial joins
- Spatial aggregation


Spatial join: approach 1

- Create index
 - Trips
- For each neighborhood
 - Query for trips within that neighborhood using the index



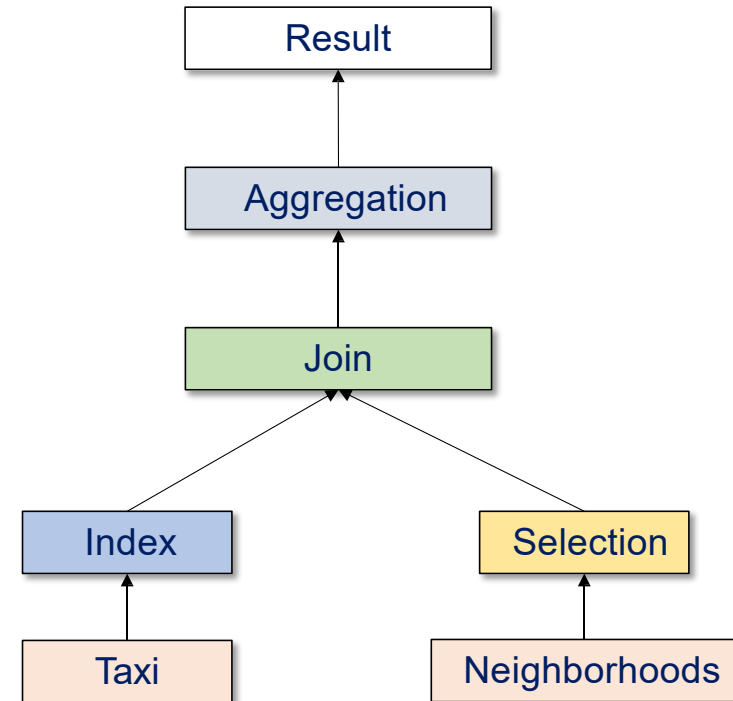
Spatial join: approach 2

- Create index
 - Neighborhoods
- For each trip
 - Query for neighborhood using the index
 - Add it to the corresponding neighborhood




Spatial aggregation queries

```
SELECT COUNT(*)  
FROM taxi T, neighborhoods N  
WHERE T.pickup inside N.polygon  
GROUP BY N.id
```

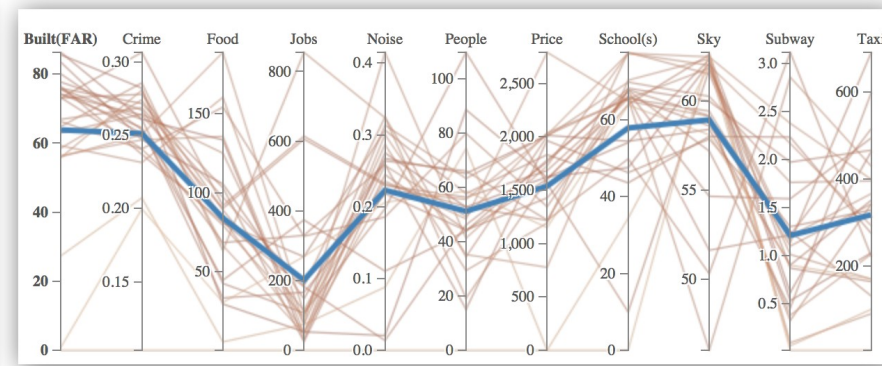




Usability through visual queries

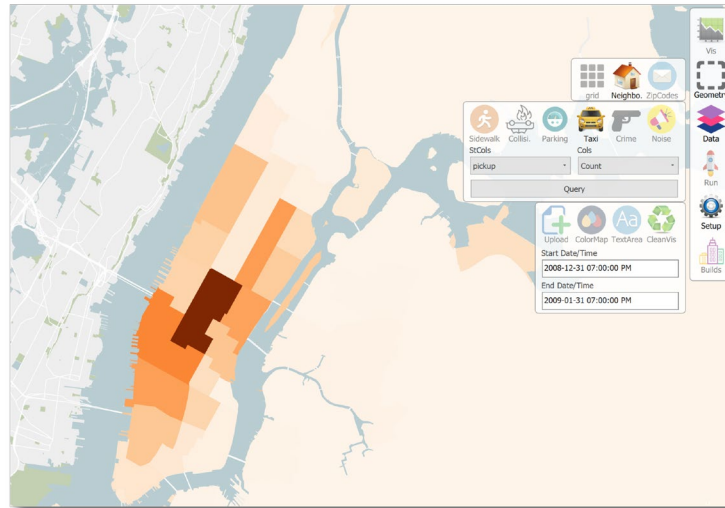


```
SELECT COUNT(*)  
FROM taxi T, neighborhoods N  
WHERE T.pickup inside N.polygon  
AND T.picktime in January 2016  
GROUP BY N.id
```

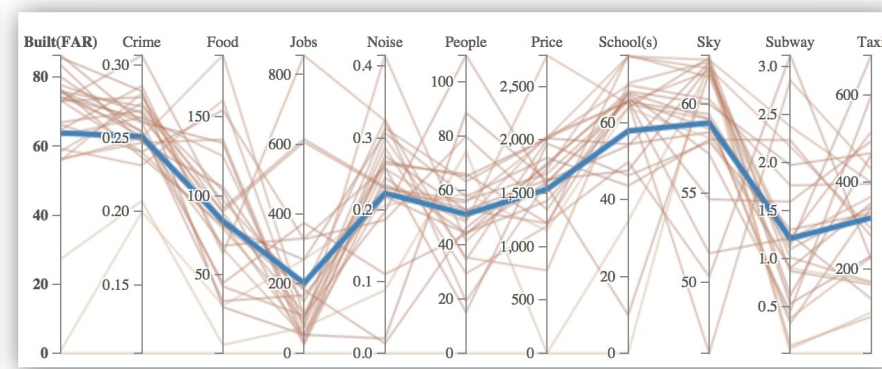


```
SELECT COUNT(*)  
FROM crime C, neighborhoods N  
WHERE C.location INSIDE N.geometry  
AND C.date in January 2016  
GROUP BY N.id
```


Usability through visual queries



```
SELECT COUNT(*)  
FROM taxi T, neighborhoods N  
WHERE T.pickup inside N.polygon  
AND T.picktime in January 2016  
GROUP BY N.id
```



```
SELECT COUNT(*)  
FROM crime C, neighborhoods N  
WHERE C.location INSIDE N.geometry  
AND C.date in January 2016  
GROUP BY N.id
```

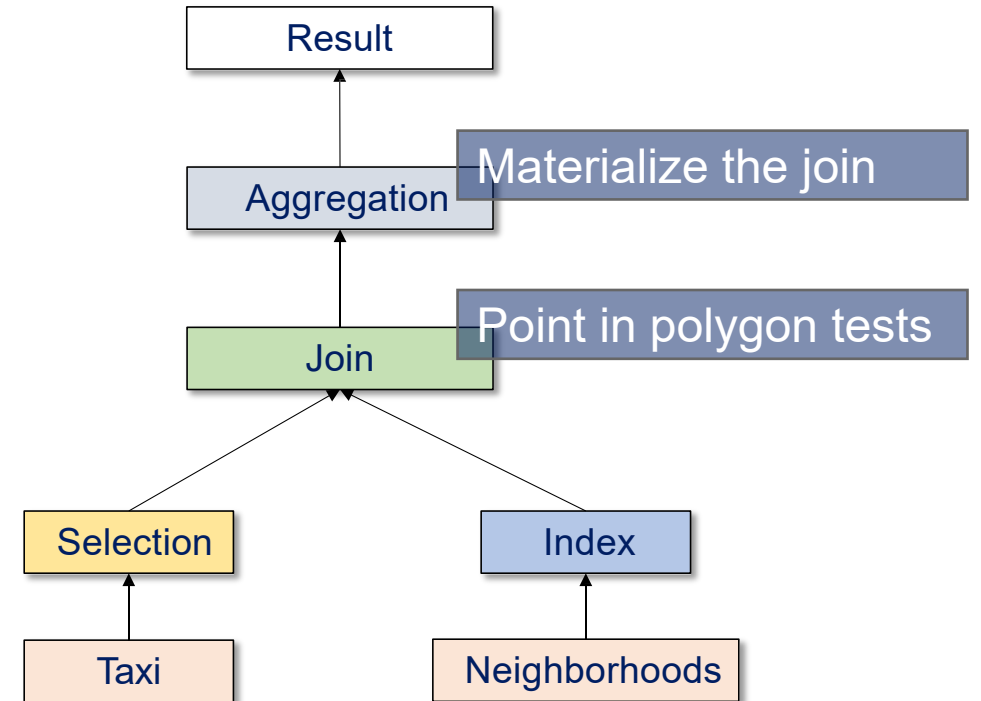
Crime
Food
Jobs
Noise
Price
Schools
Sky
exposure

Spatial aggregation queries

- Set of trips
 - Point data
 - ~340 million trips
- Set of neighborhoods
 - Polygon data
 - ~260 neighborhoods
- How to join these two data sets?

Spatial aggregation queries

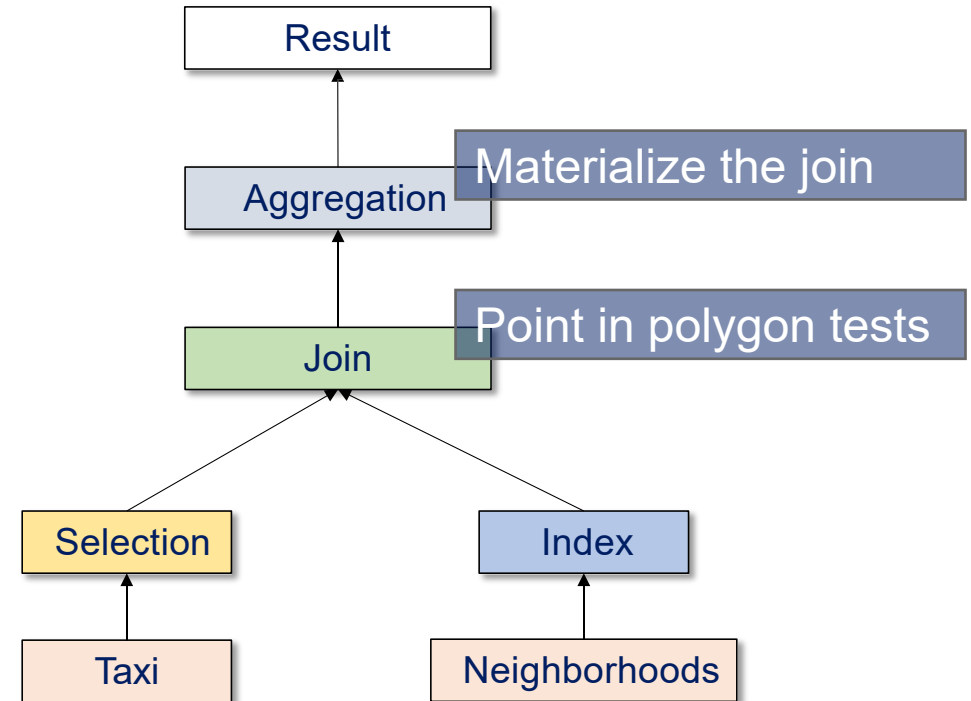
```
SELECT COUNT(*)  
FROM taxi T, neighborhoods N  
WHERE T.pickup inside N.polygon  
GROUP BY N.id
```



Spatial aggregation queries

```
SELECT COUNT(*)  
FROM taxi T, neighborhoods N  
WHERE T.pickup inside N.polygon  
GROUP BY N.id
```


- Existing spatial databases.
- Several minutes!



Spatial aggregation queries

```
SELECT COUNT(*)  
FROM taxi T, neighborhoods N  
WHERE T.pickup inside N.polygon  
GROUP BY N.id
```


- Possible solution: pre-compute aggregation



Spatial aggregation queries

```
SELECT COUNT(*)  
FROM taxi T, neighborhoods N  
WHERE T.pickup inside N.polygon  
AND T.picktime in March 2011  
GROUP BY N.id
```


- Possible solution: cube-based structures – nanocubes, ...



Spatial aggregation queries


```
SELECT COUNT(*)  
FROM taxi T, neighborhoods N  
WHERE T.pickup inside N.polygon  
AND T.picktime in March 2011  
AND T.duration > 10 minutes  
GROUP BY N.id
```

- Possible solution: cube-based structures – nanocubes, ...
- Space explosion!





Desiderata

- Interactive response times
- Avoid costly preprocessing




Running example

```
SELECT COUNT(*)  
FROM taxi T, neighborhoods N  
WHERE T.pickup inside N.polygon  
GROUP BY N.id
```




Running example

```
SELECT COUNT(*)  
FROM taxi T, neighborhoods N  
WHERE T.pickup inside N.polygon  
GROUP BY N.id
```




Spatial aggregation


- A geometric perspective of spatial aggregation.
- Spatial join: “drawing” points and polygons on the same canvas.



Input points



Input polygon




Spatial join

- Leverage the graphics pipeline of the GPU [Tzirita et al., 2017]


Raster Join

```
SELECT COUNT(*)  
FROM taxi T, neighborhoods N  
WHERE T.pickup inside N.polygon  
GROUP BY N.id
```




Raster Join

```
SELECT COUNT(*)  
FROM taxi T, neighborhoods N  
WHERE T.pickup inside N.polygon  
GROUP BY N.id
```




Raster Join

```
SELECT COUNT(*)  
FROM taxi T, neighborhoods N  
WHERE T.pickup inside N.polygon  
GROUP BY N.id
```




Raster Join

```
SELECT COUNT(*)  
FROM taxi T, neighborhoods N  
WHERE T.pickup inside N.polygon  
GROUP BY N.id
```




Raster Join

```
SELECT COUNT(*)  
FROM taxi T, neighborhoods N  
WHERE T.pickup inside N.polygon  
GROUP BY N.id
```




Raster Join

```
SELECT COUNT(*)  
FROM taxi T, neighborhoods N  
WHERE T.pickup inside N.polygon  
GROUP BY N.id
```




Raster Join

```
SELECT COUNT(*)  
FROM taxi T, neighborhoods N  
WHERE T.pickup inside N.polygon  
GROUP BY N.id
```




Raster Join

```
SELECT COUNT(*)  
FROM taxi T, neighborhoods N  
WHERE T.pickup inside N.polygon  
GROUP BY N.id
```



Raster Join

```
SELECT COUNT(*)  
FROM taxi T, neighborhoods N  
WHERE T.pickup inside N.polygon  
GROUP BY N.id
```




7

[illegible]



Raster Join

- Exploits native support for drawing in GPUs.
- Combines the aggregation with the join operation.
- No Point-in-Polygon tests.
- Shortcomings?



Raster Join: approximation


- Rasterization can introduce false positives and false negatives.
- Errors can be reduced: approximate the polygon outline by increasing screen resolution (i.e., reducing pixel size).
- Accurate Raster Join: point-in-polygon tests for points in the boundary.



Raster Join: performance evaluation

NYC taxi data (over 868 million points), 260 NYC neighborhood polygons

Laptop with i7 Quad-Core@2.8 GHz, 16 GB RAM, GTX 1060 GPU



300x speedup

